

LIBP2P: A MODULAR PEER-TO-PEER NETWORKING STACK

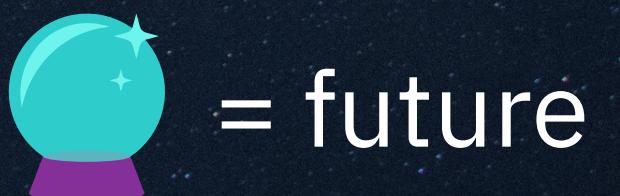
Crosslink Taipei 2019 

Raúl Kripalani

tech lead, libp2p @ Protocol Labs



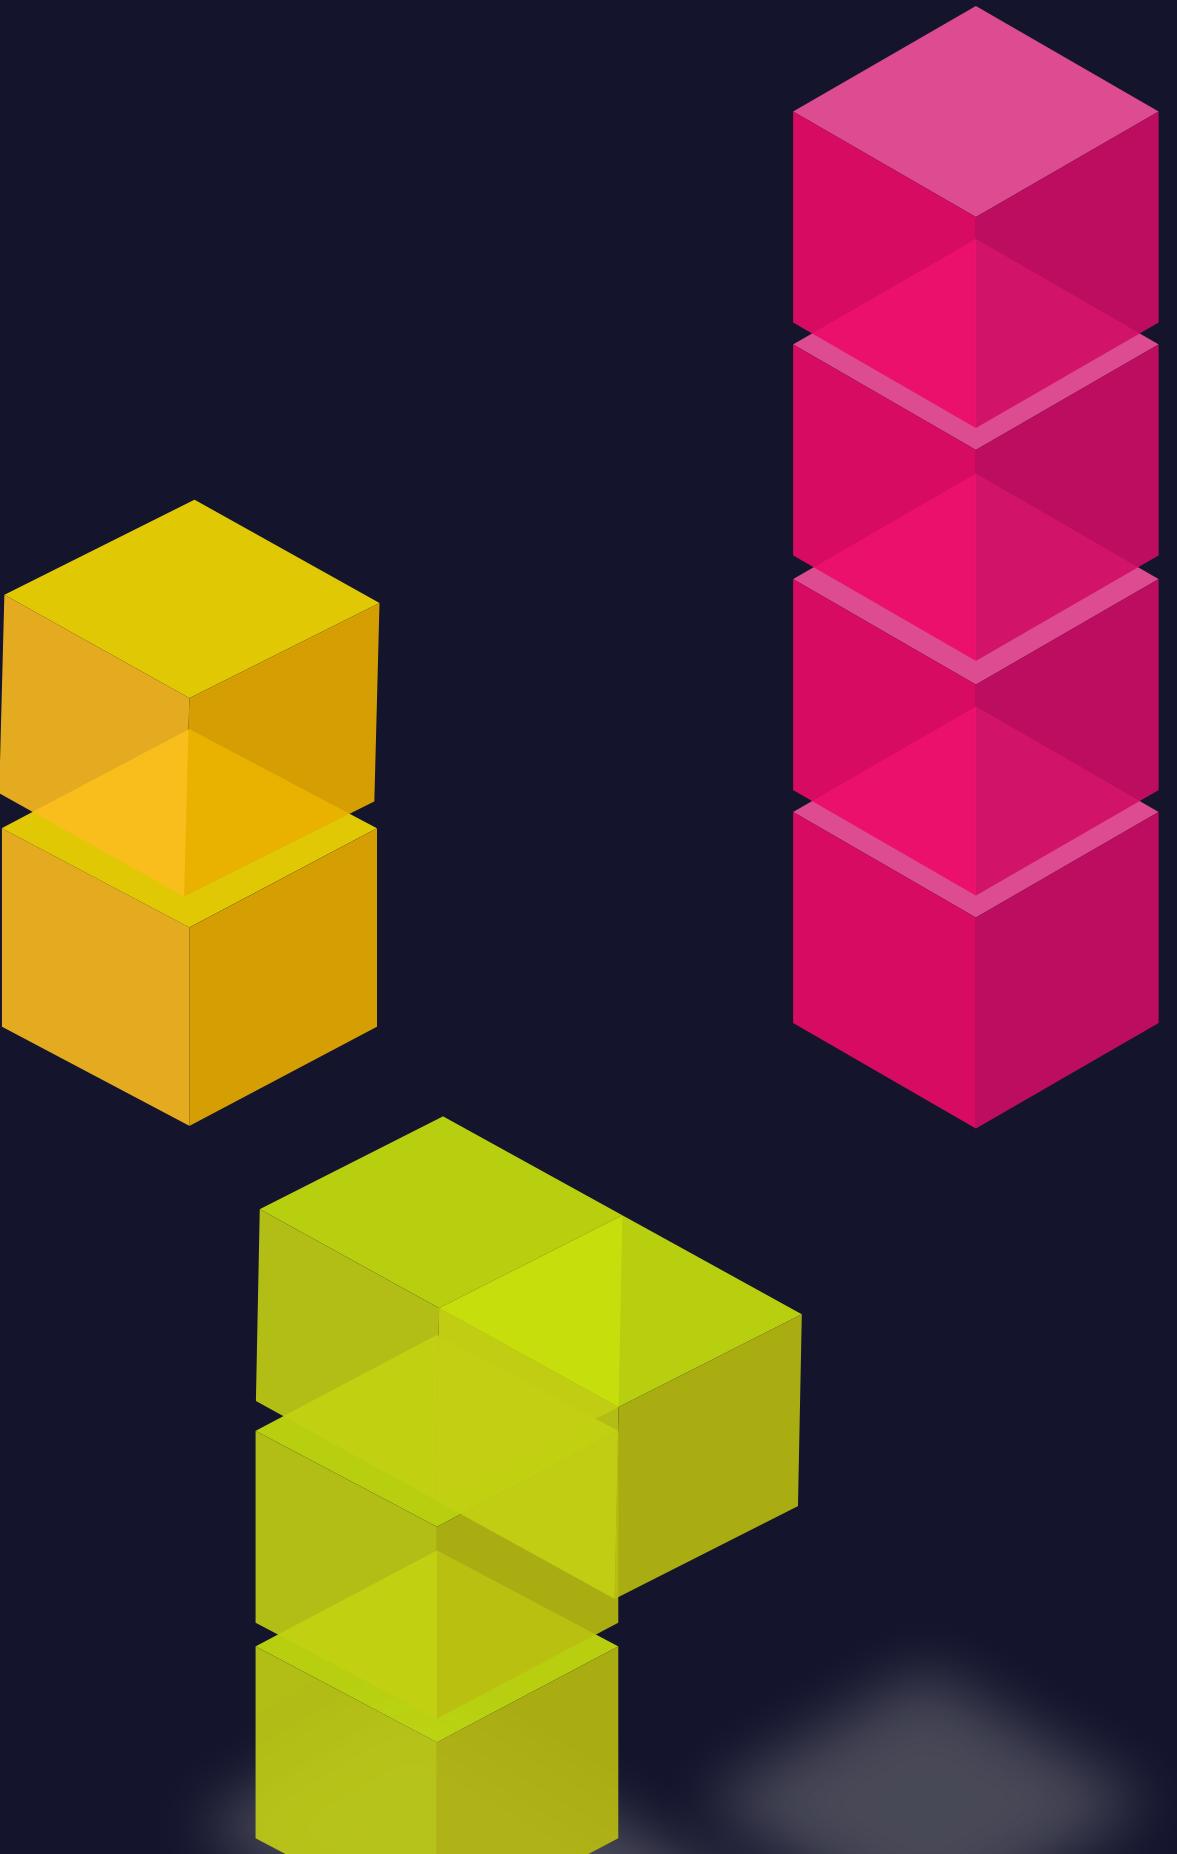
libp2p



today

AGENDA

- 
- introduction
 - addressing and identity
 - transports
 - connection bootstrapping
 - secure channels
 - multiplexing
 - protocol negotiation
 - peer discovery and content routing
 - pubsub
 - connectivity (NAT, relays)
 - what's next?





A MODULAR PEER-TO-PEER NETWORKING STACK

Composable building blocks to assemble ***future-proof p2p networking layers.***

Runs on many runtimes: server, browser, mobile (experiments). soon: embedded.

Originated in IPFS. Implemented in 7 languages; bustling cross-project community.

Created, stewarded and sponsored by PL; owned by the community. It's a public good.

Licensed under MIT, soon Permissive License Stack (ASLv2 + MIT).



Transports

Multiplexers

Secure Channels

Peer Discovery

Pubsub

NAT Traversal

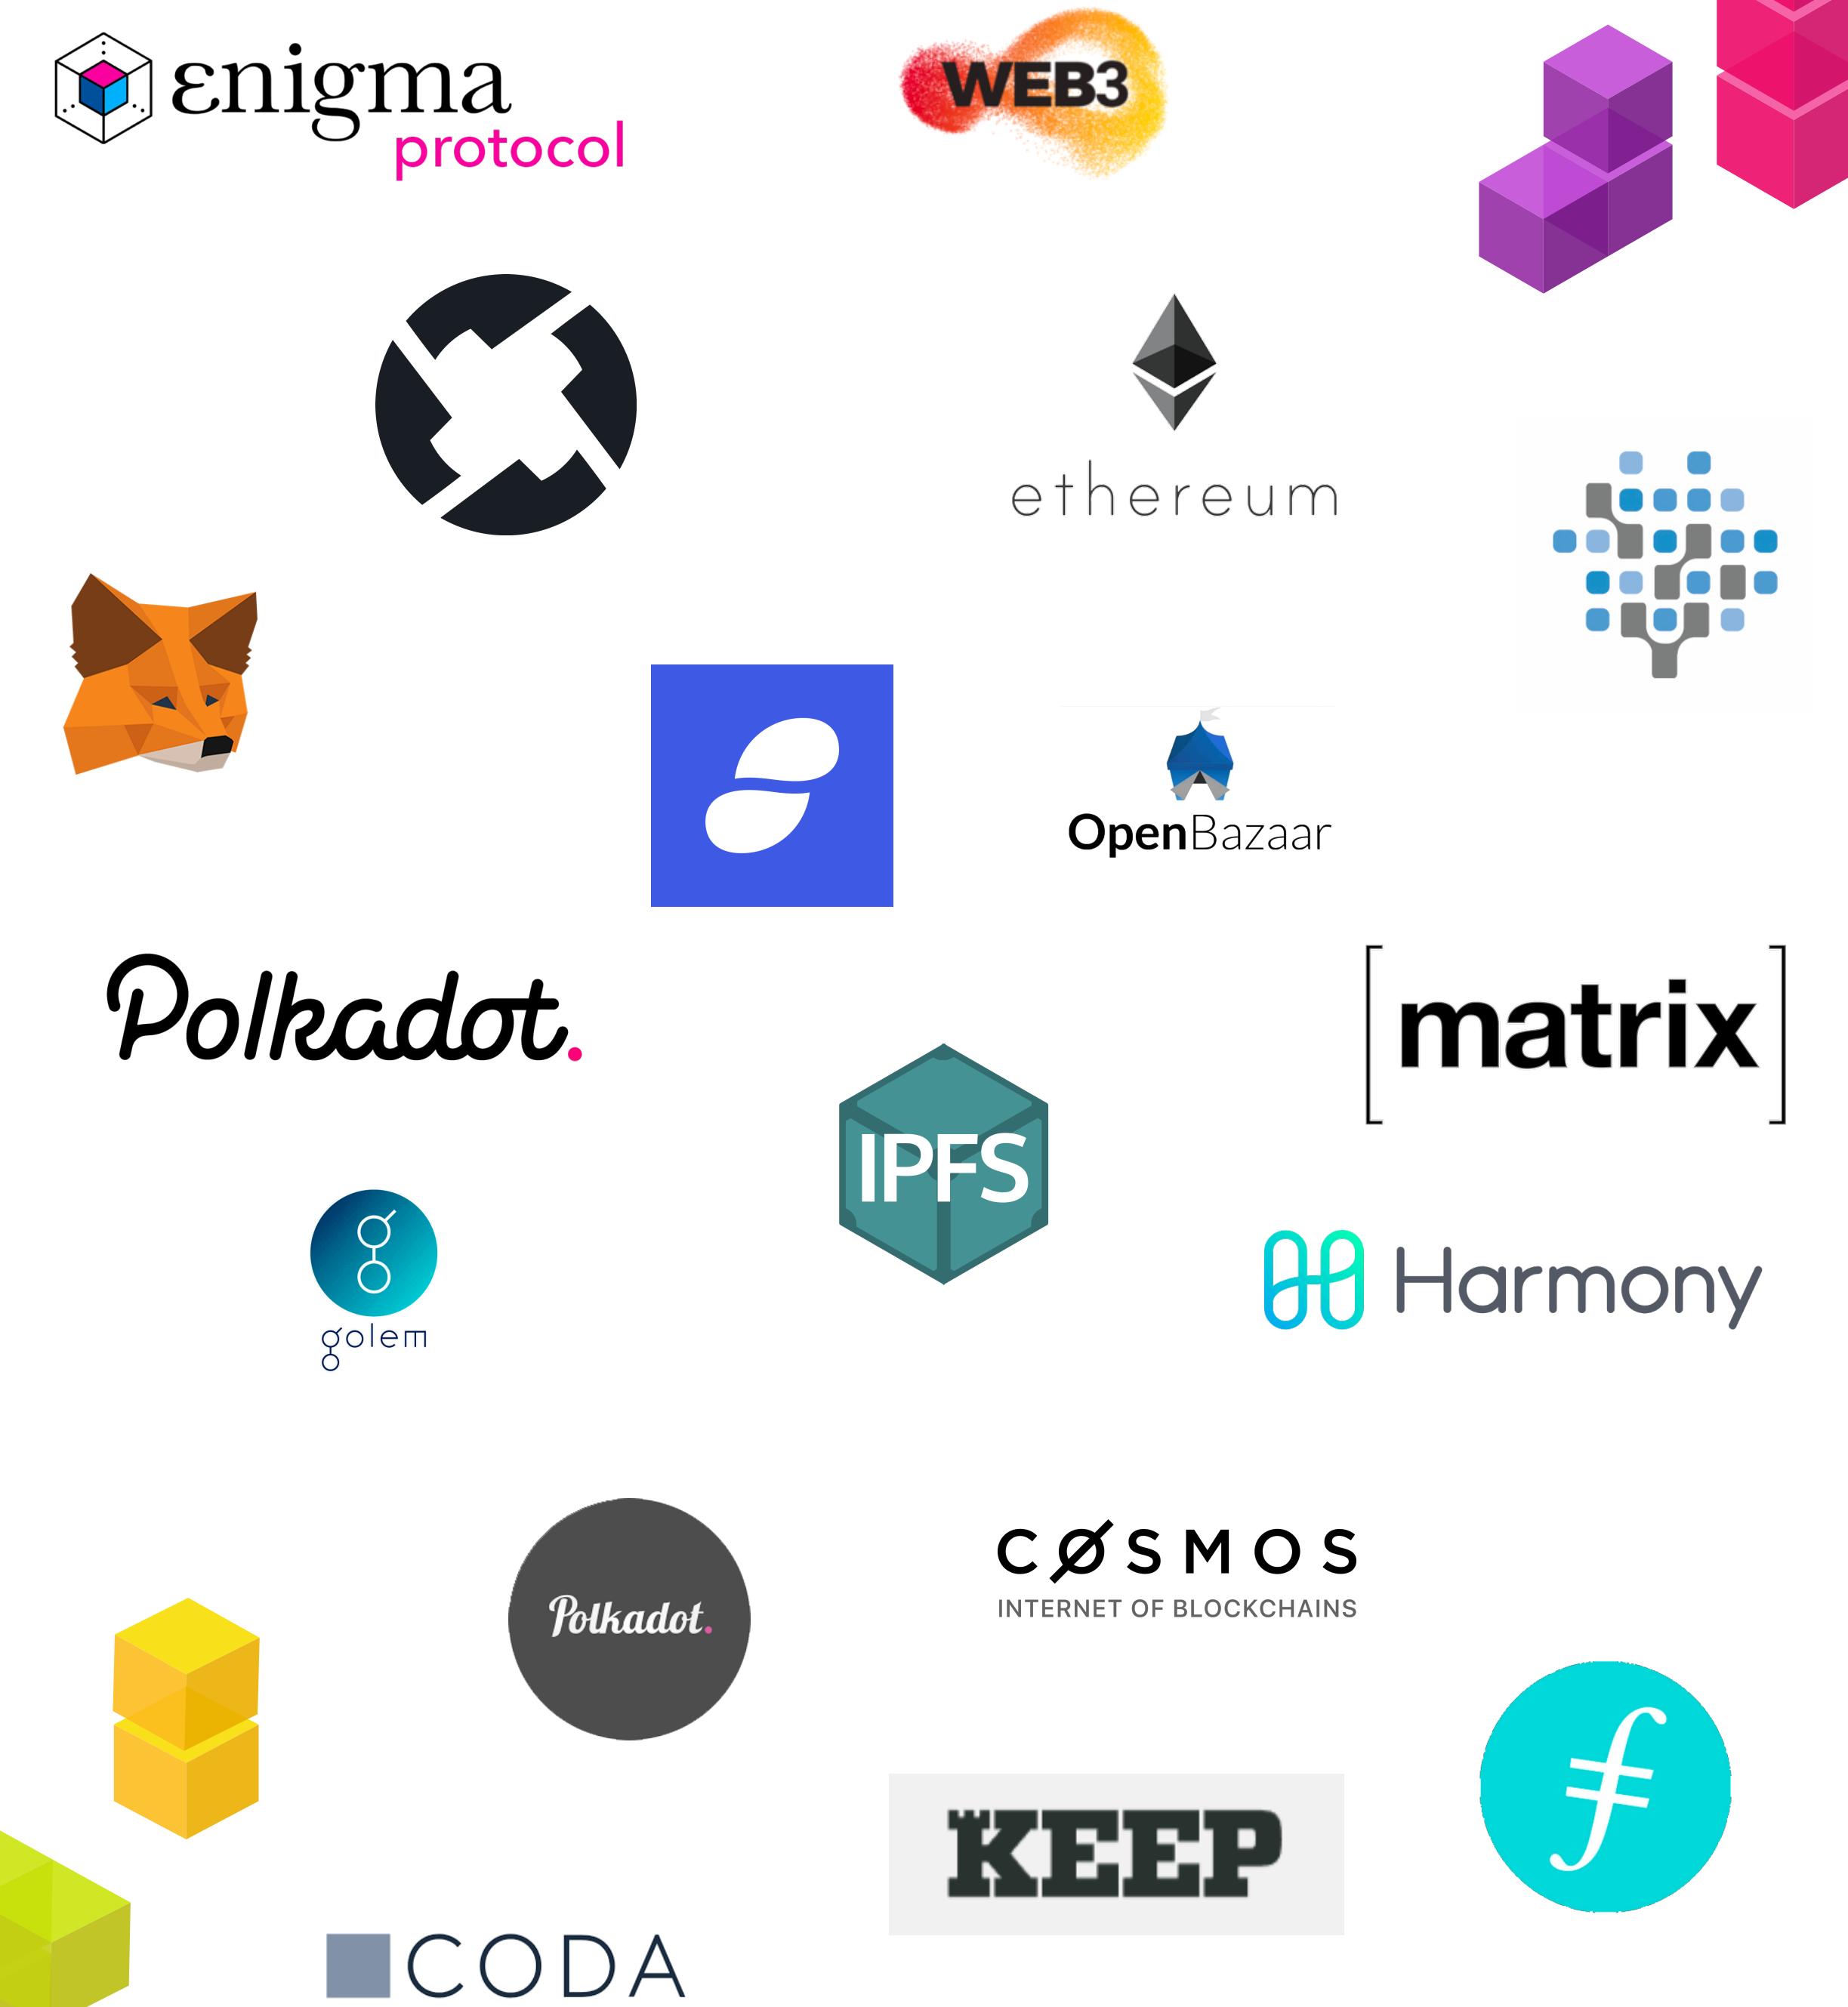
Peer Routing

Content Routing

libp2p

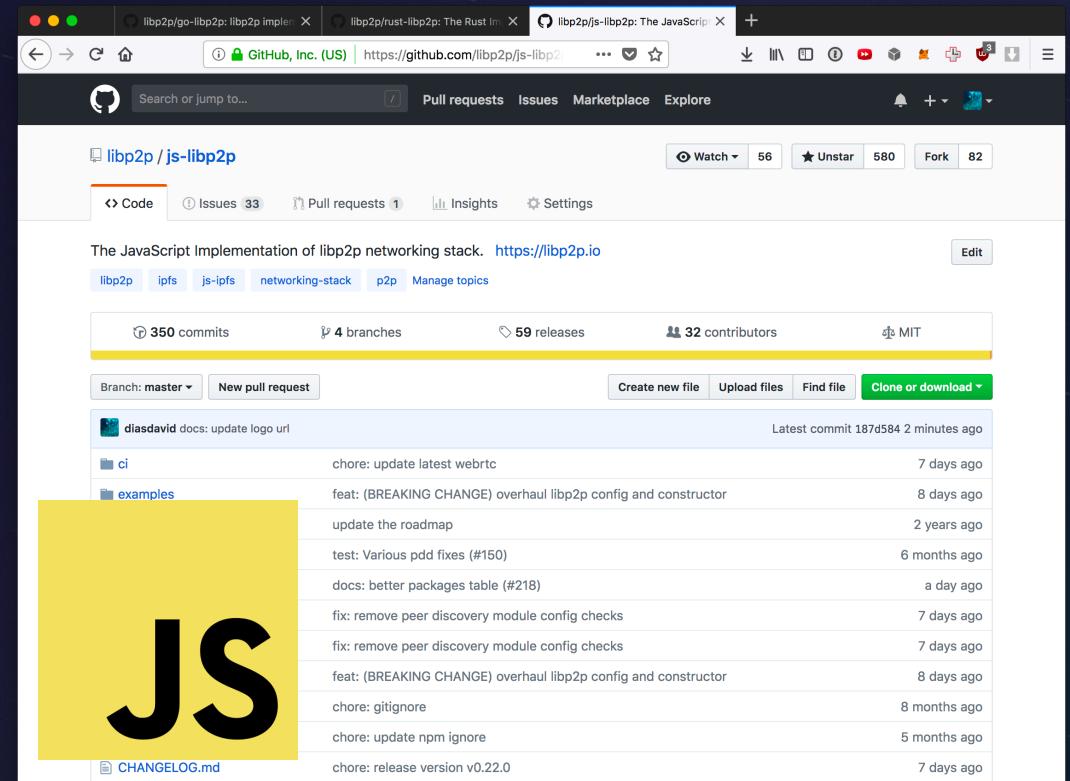
ADOPTERS

IPFS
Polkadot
Ethereum 2.0
Filecoin
MetaMask
OpenBazaar
0x
Golem
Harmony
Keep Network
Validity Labs
Enigma
Cosmos (considering)
Coda Protocol (considering)
Matrix (experimenting)
... more



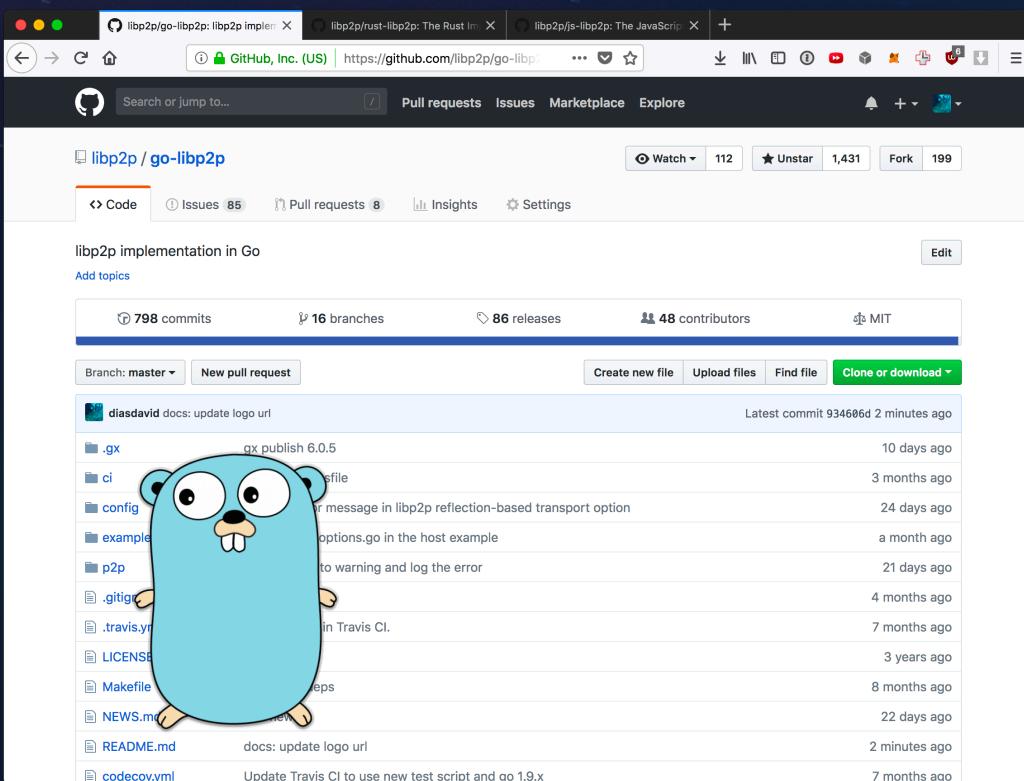
7 native implementations

js-libp2p



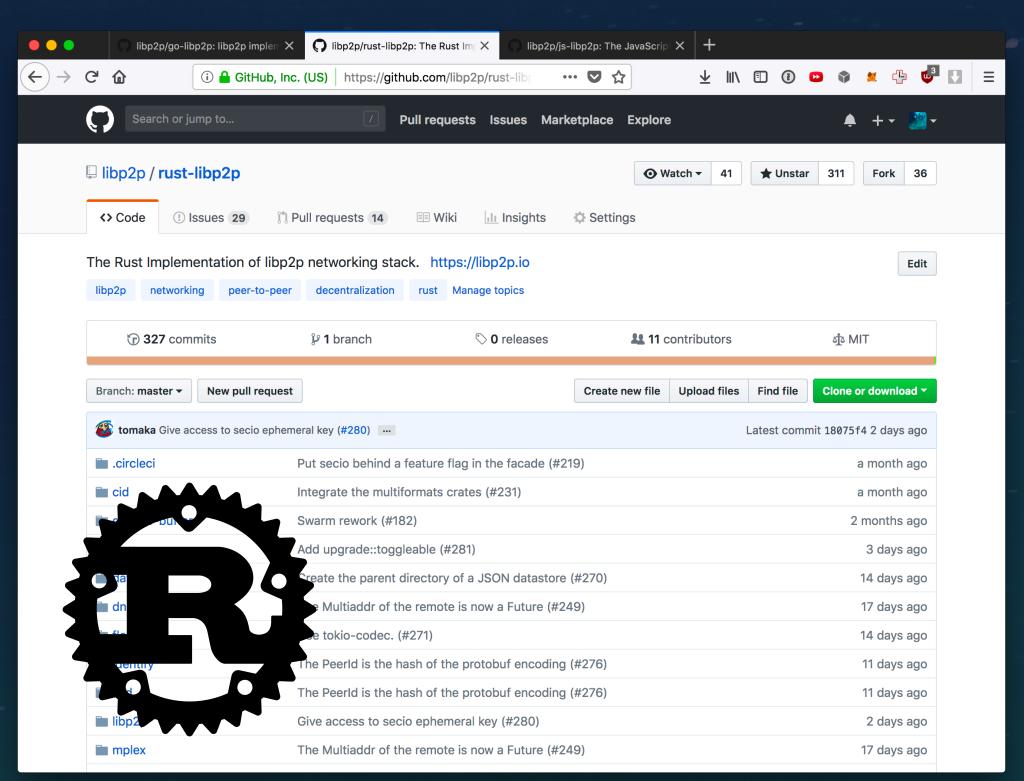
Protocol Labs

go-libp2p



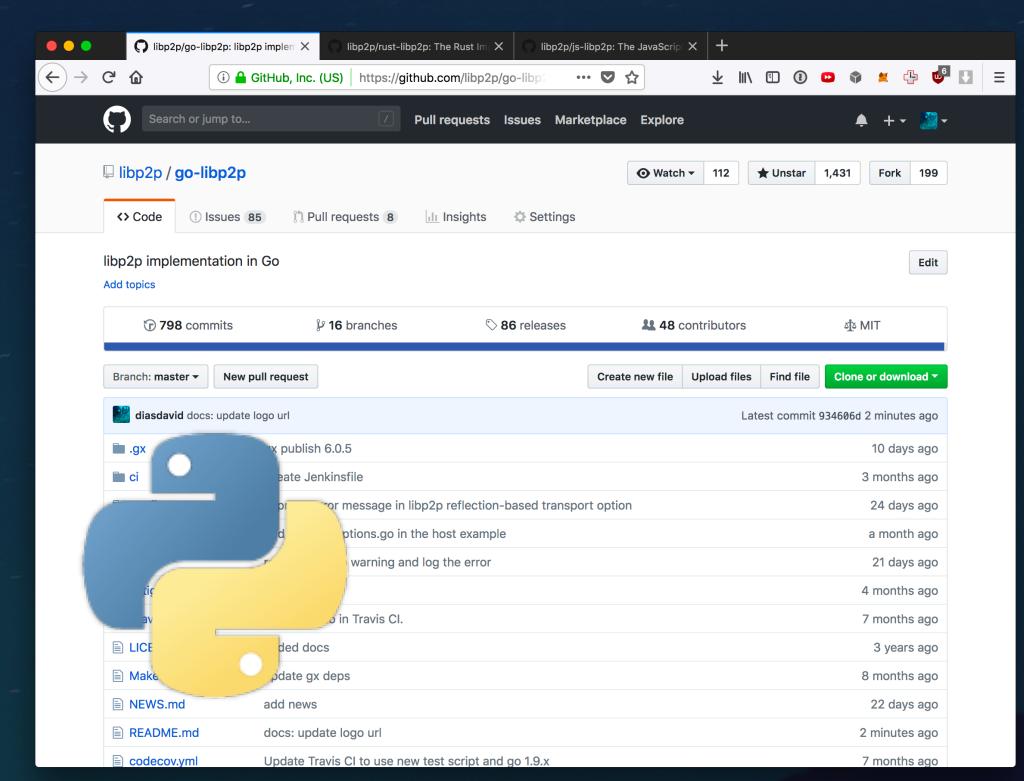
Protocol Labs

rust-libp2p



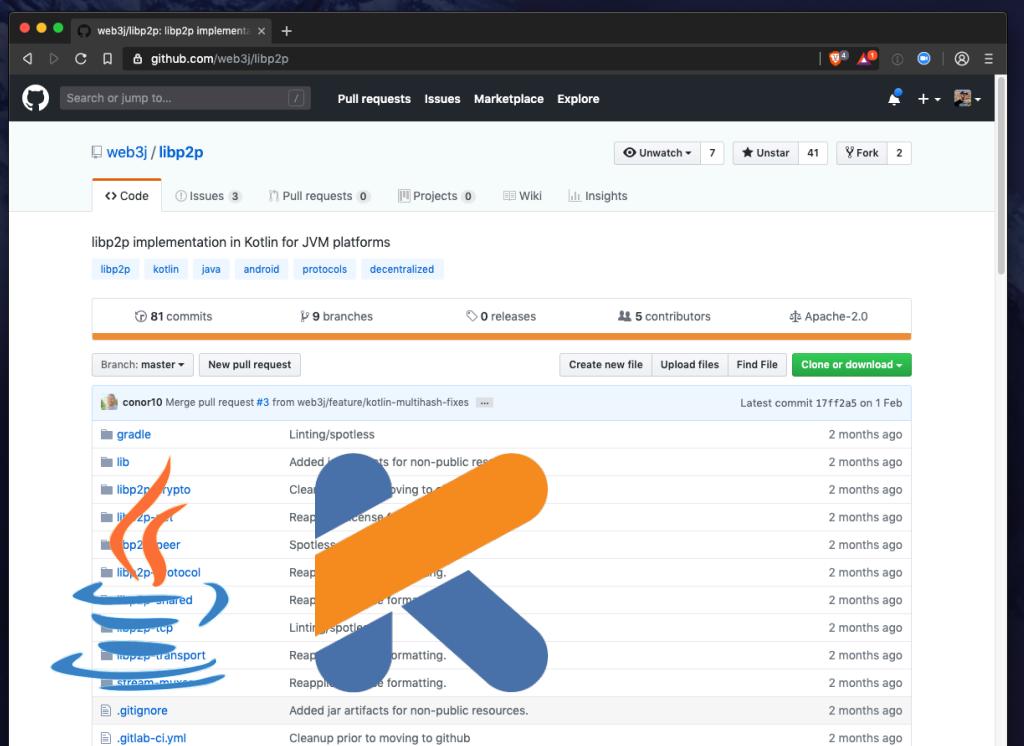
Web3 Foundation

py-libp2p



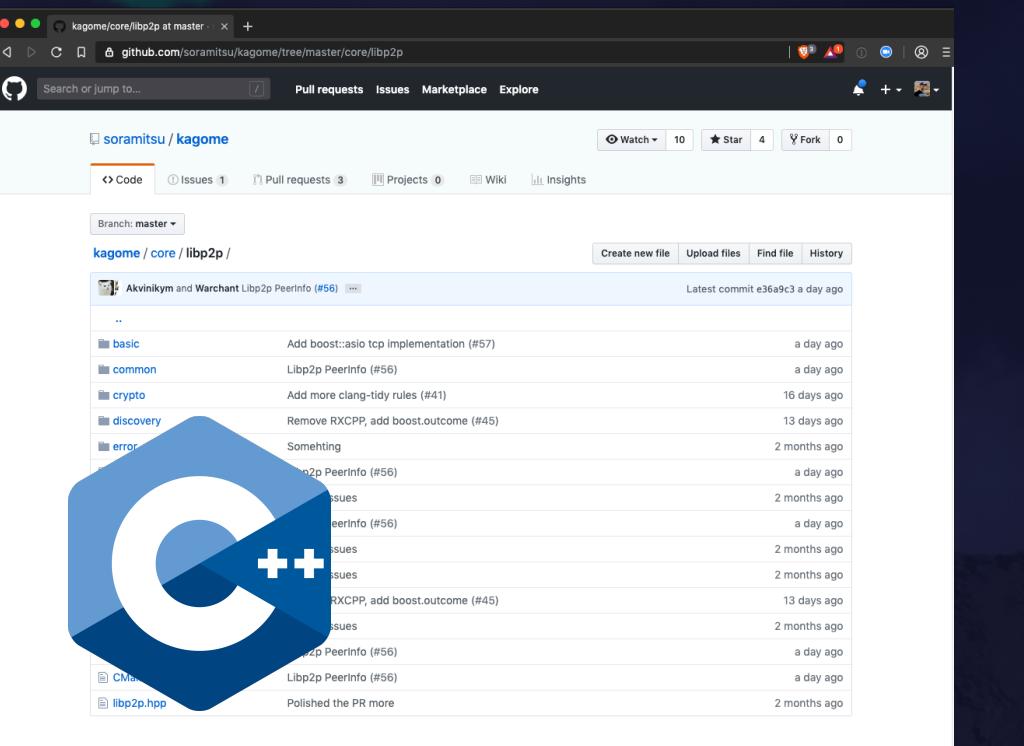
Ethereum Foundation

jvm-libp2p



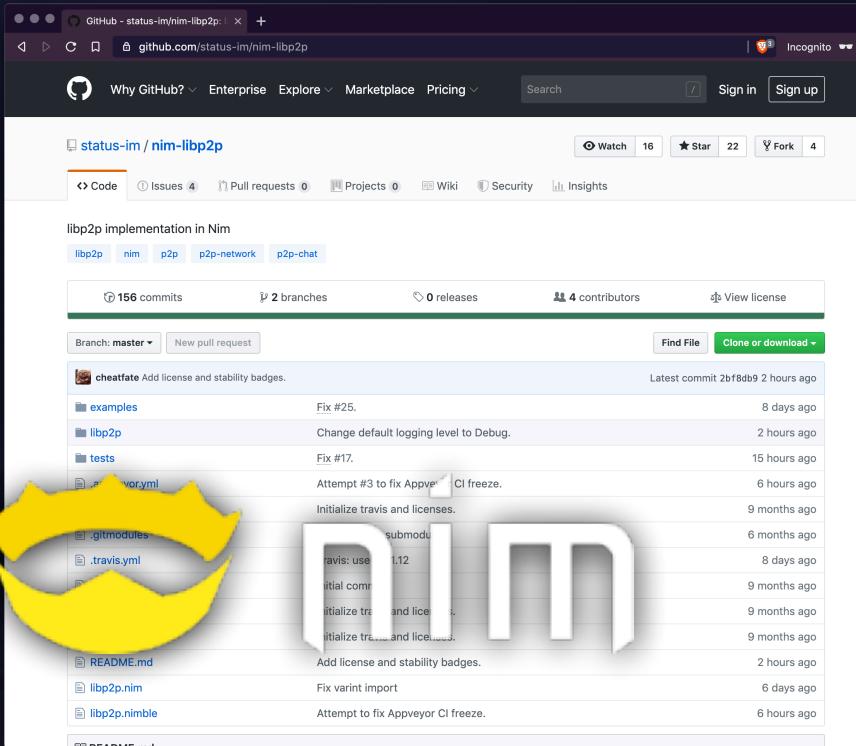
Harmony, PegaSys, Web3 Labs

cpp-libp2p



Web3 Foundation, Soramitsu

nim-libp2p



Status



ETH2 NETWORKING



Ethereum 2.0 networking specification

Table of contents

- Network fundamentals
 - Transport
 - Encryption and identification
 - Protocol negotiation
 - Multiplexing
- Eth 2.0 network interaction domains
 - Configuration
 - The gossip domain: gossipsub
 - The Req/Resp domain
 - The discovery domain: discv5
- Design decision rationale
 - Transport
 - Multiplexing
 - Protocol negotiation
 - Encryption
 - Gossipsub
 - Req/Resp
 - Discovery
 - Compression/Encoding
- libp2p implementations matrix



DECENTRALIZED PROCESS ADDRESSING

libp2p's raison d'être is the ability to **locate**, **connect**, **authenticate**, **negotiate** and **interact** efficiently with any process in the world, ***no matter the runtime*** (server, browser, IoT, embedded, etc.) so long as its ***identity is cryptographically derived from its public key***; and have all of that happen in a ***seamless manner*** (e.g. NAT, relay, packet switching), even as those processes **relocate**, **roam**, **evolve** and **mutate over time**. It is ***juxtaposed to endpoint addressing*** (e.g. IP networks).





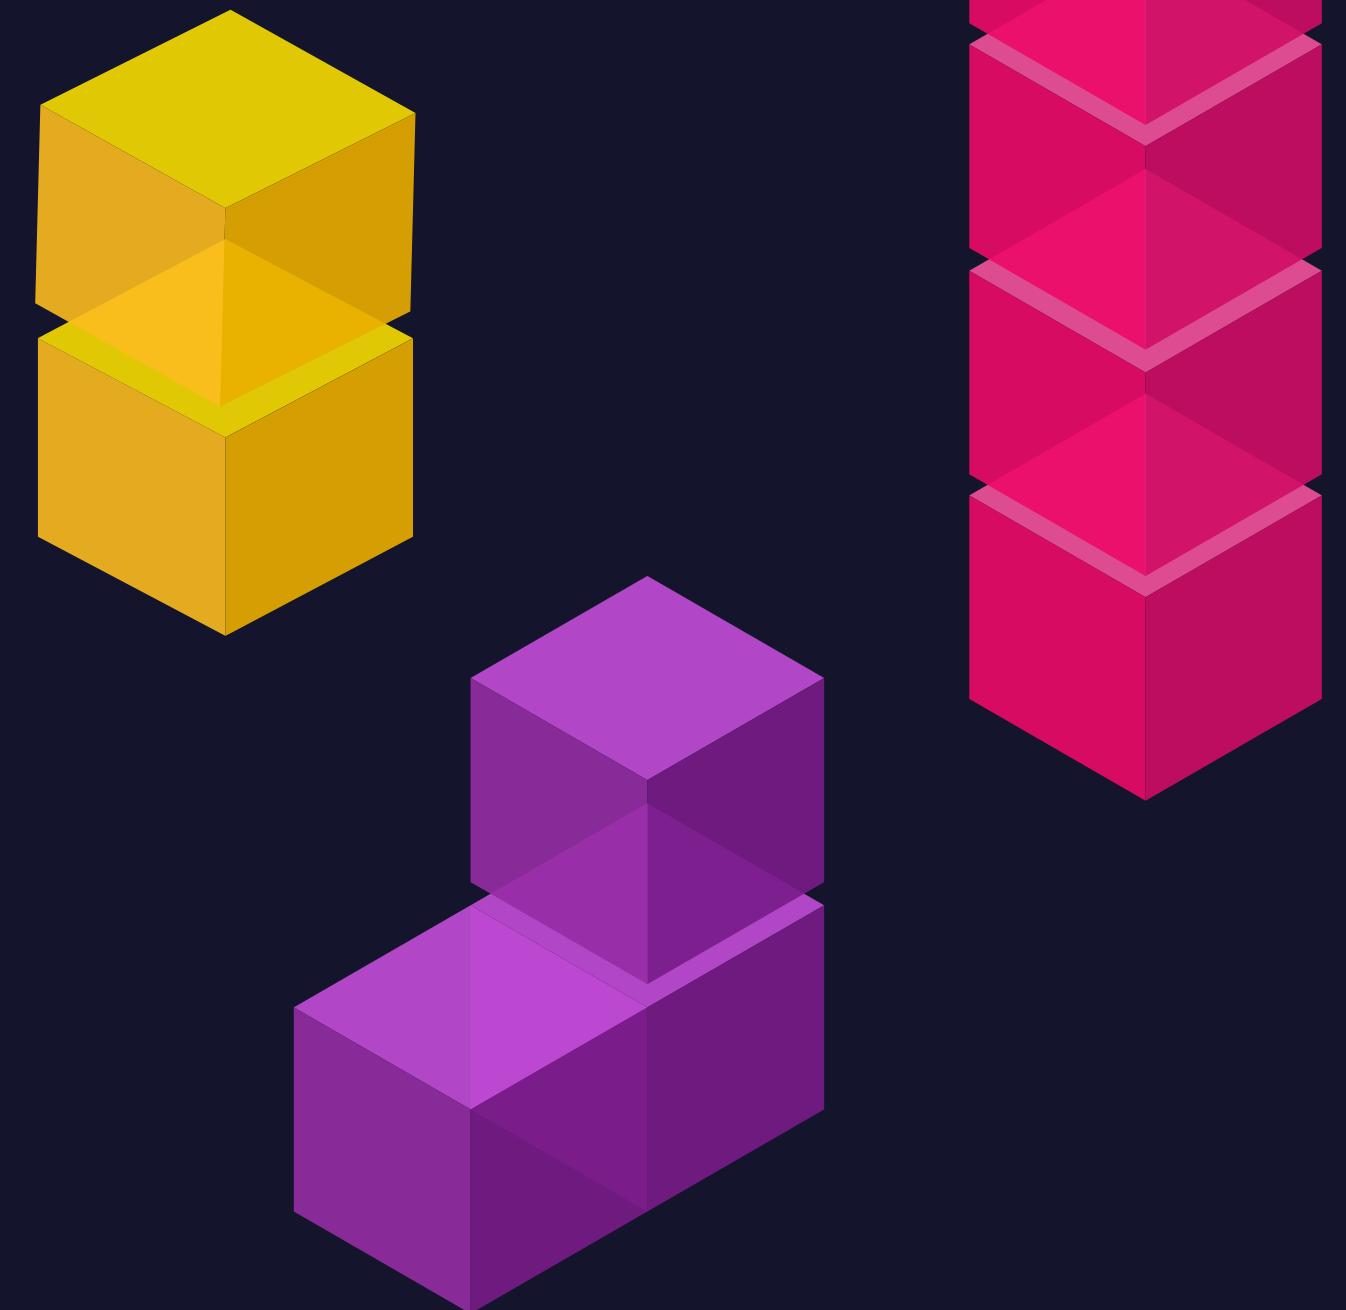
addressing and identity

transport-agnostic addresses

MULTIADDRS

/ip4/104.236.179.241/tcp/4001/p2p/QmPeer...

0x0468ecb3f1060fa1 (omitting the id)



- Composable, future-proof, upgradeable transport-agnostic addresses.
- binary byte-encoded packed format, with a canonical string representation.
- encodes addressing, transport, routing, identity, encryption (future) of a peer.
- flexible varint-based byte[] format vs. hardcoding assumptions.
- what if...?

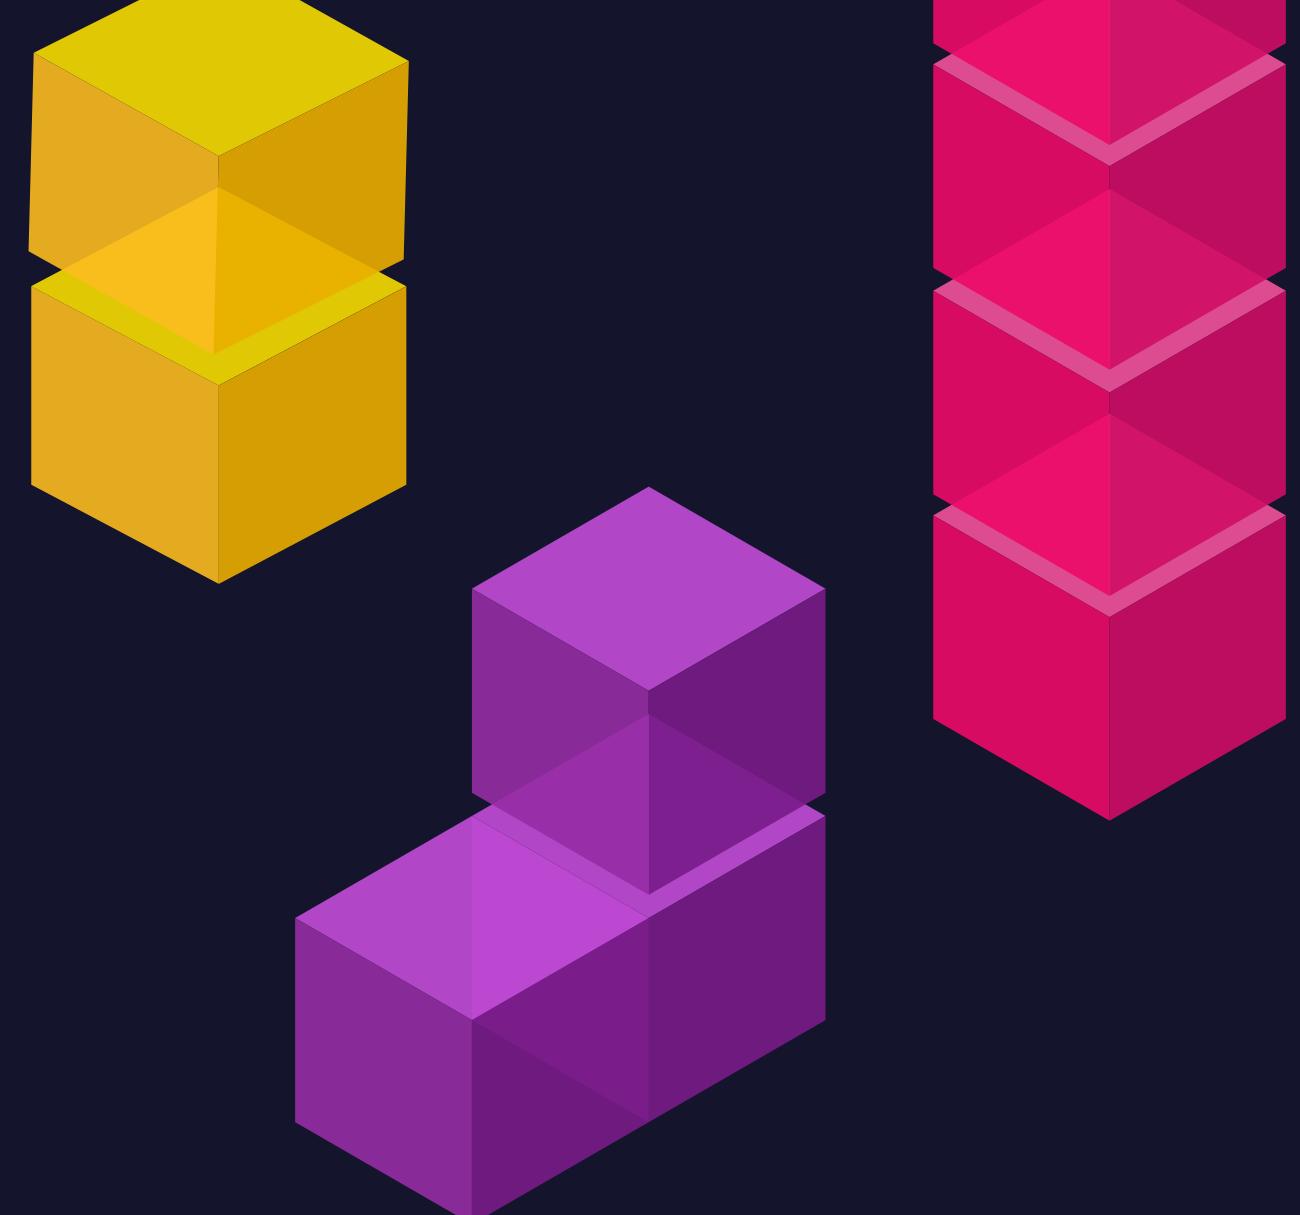


addressing and identity

multiaddrs are inspired by the

PLAN9 UBIQUITOUS FILESYSTEM

everything is a file, including remote network processes



```
/  
|_ ip4          # all processes addressable via IPv4  
    |_ 7.7.7.7    # the process lives in network address 7.7.7.7  
        |_ udp     # all processes exposing udp sockets  
        |_ tcp     # all processes exposing tcp sockets  
            |_ 1111  # the process sitting in port 1111  
            |_ 1234  # the process sitting in port 1234  
                |_ p2p   # all identities exposed by this process  
                    |_ QmPeerA # we want to talk to identity QmPeerA  
                    |_ QmPeerB # we want to talk to identity QmPeerB  
                        |_ p2p-circuit # relay  
                            |_ ... you get it ...
```

http://doc.cat-v.org/plan_9/misc/ubiquitous_fileserver/



PEER IDENTITY

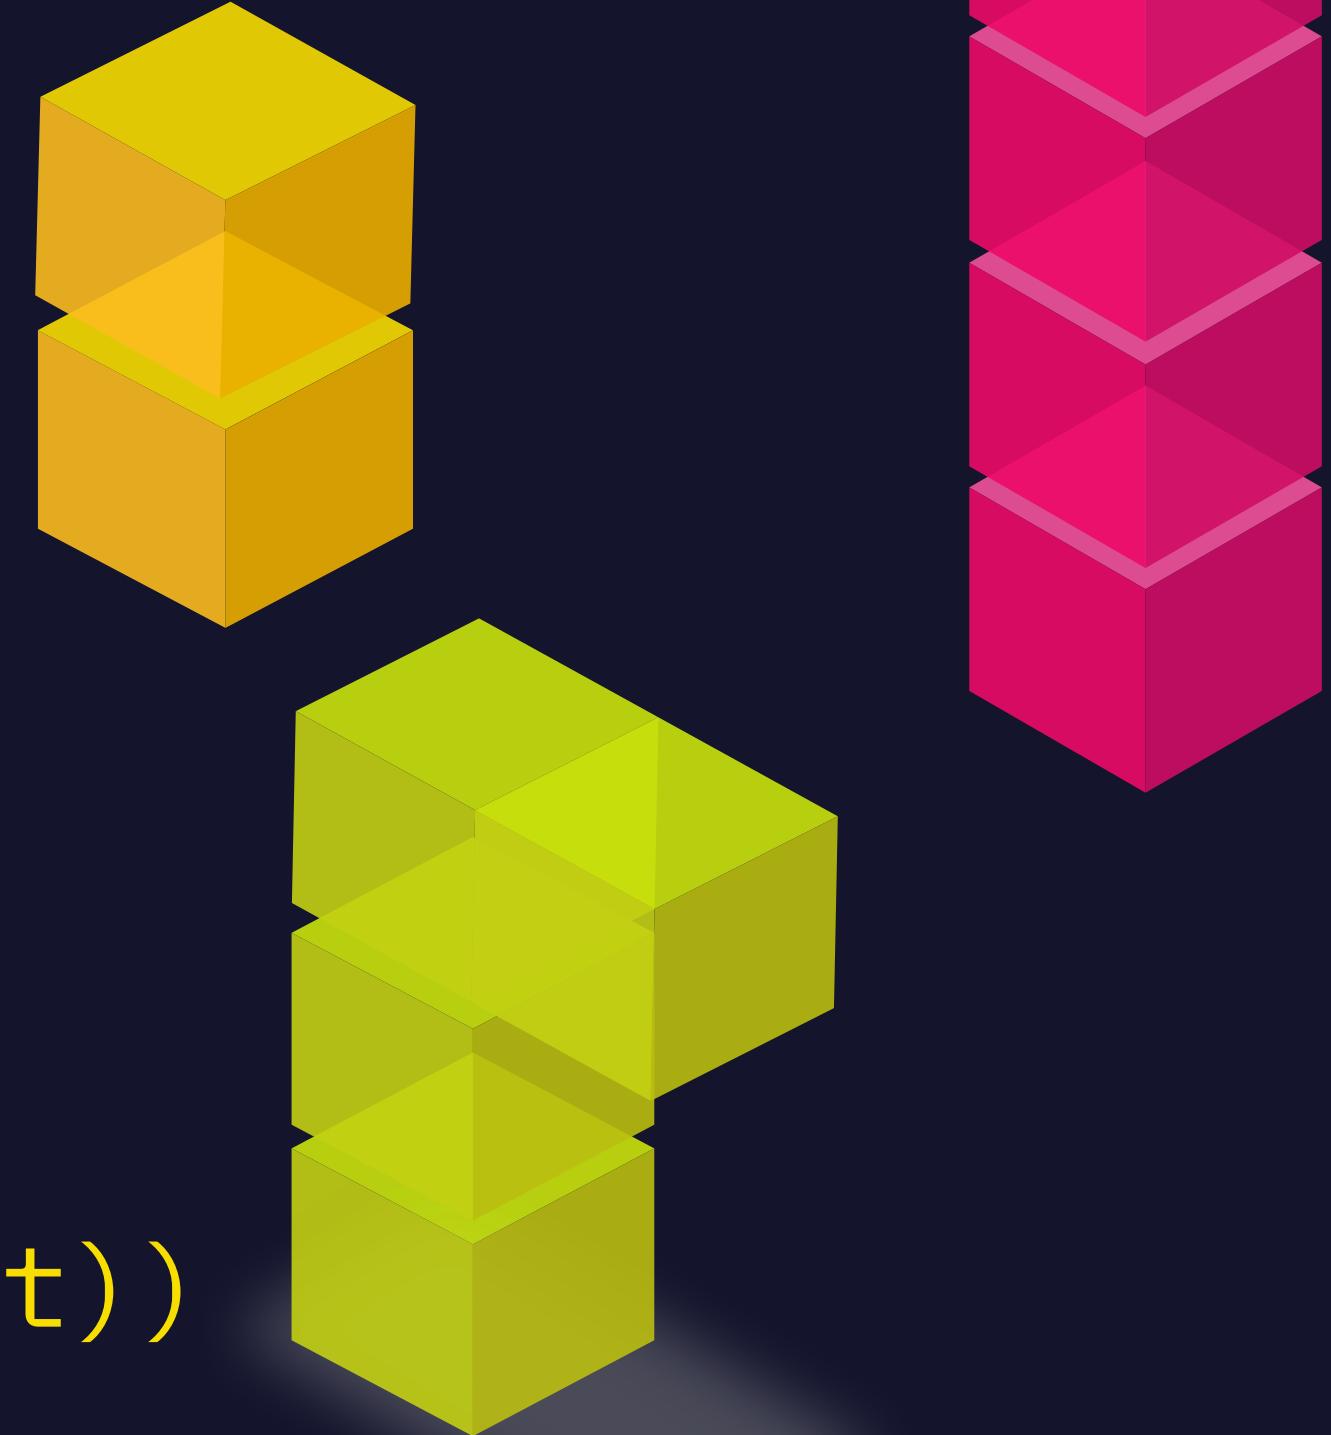
QmSoLPppuBtQSGwKDZT2M73ULpjvfd3aZ6ha4oFGL1KrGM

pubkey-struct := { <key-type>, <pubkey-bytes> }

peer-id := multihash(SHA-256, protobuf(pubkey-struct))

- Peer IDs: derived from the public key.
- Four key types: RSA 🤢, Ed25519 🤘, SECP256k1, ECDSA.
- Serialize into pb, then multihash with SHA-256.
- String representation: base58btc.
- multiaddrs: /ip4/1.2.3.4/tcp/5001/p2p/<peer-id>

```
enum KeyType {  
    RSA = 0;  
    Ed25519 = 1;  
    Secp256k1 = 2;  
    ECDSA = 3;  
}  
  
message PublicKey {  
    required KeyType Type = 1;  
    required bytes Data = 2;  
}
```

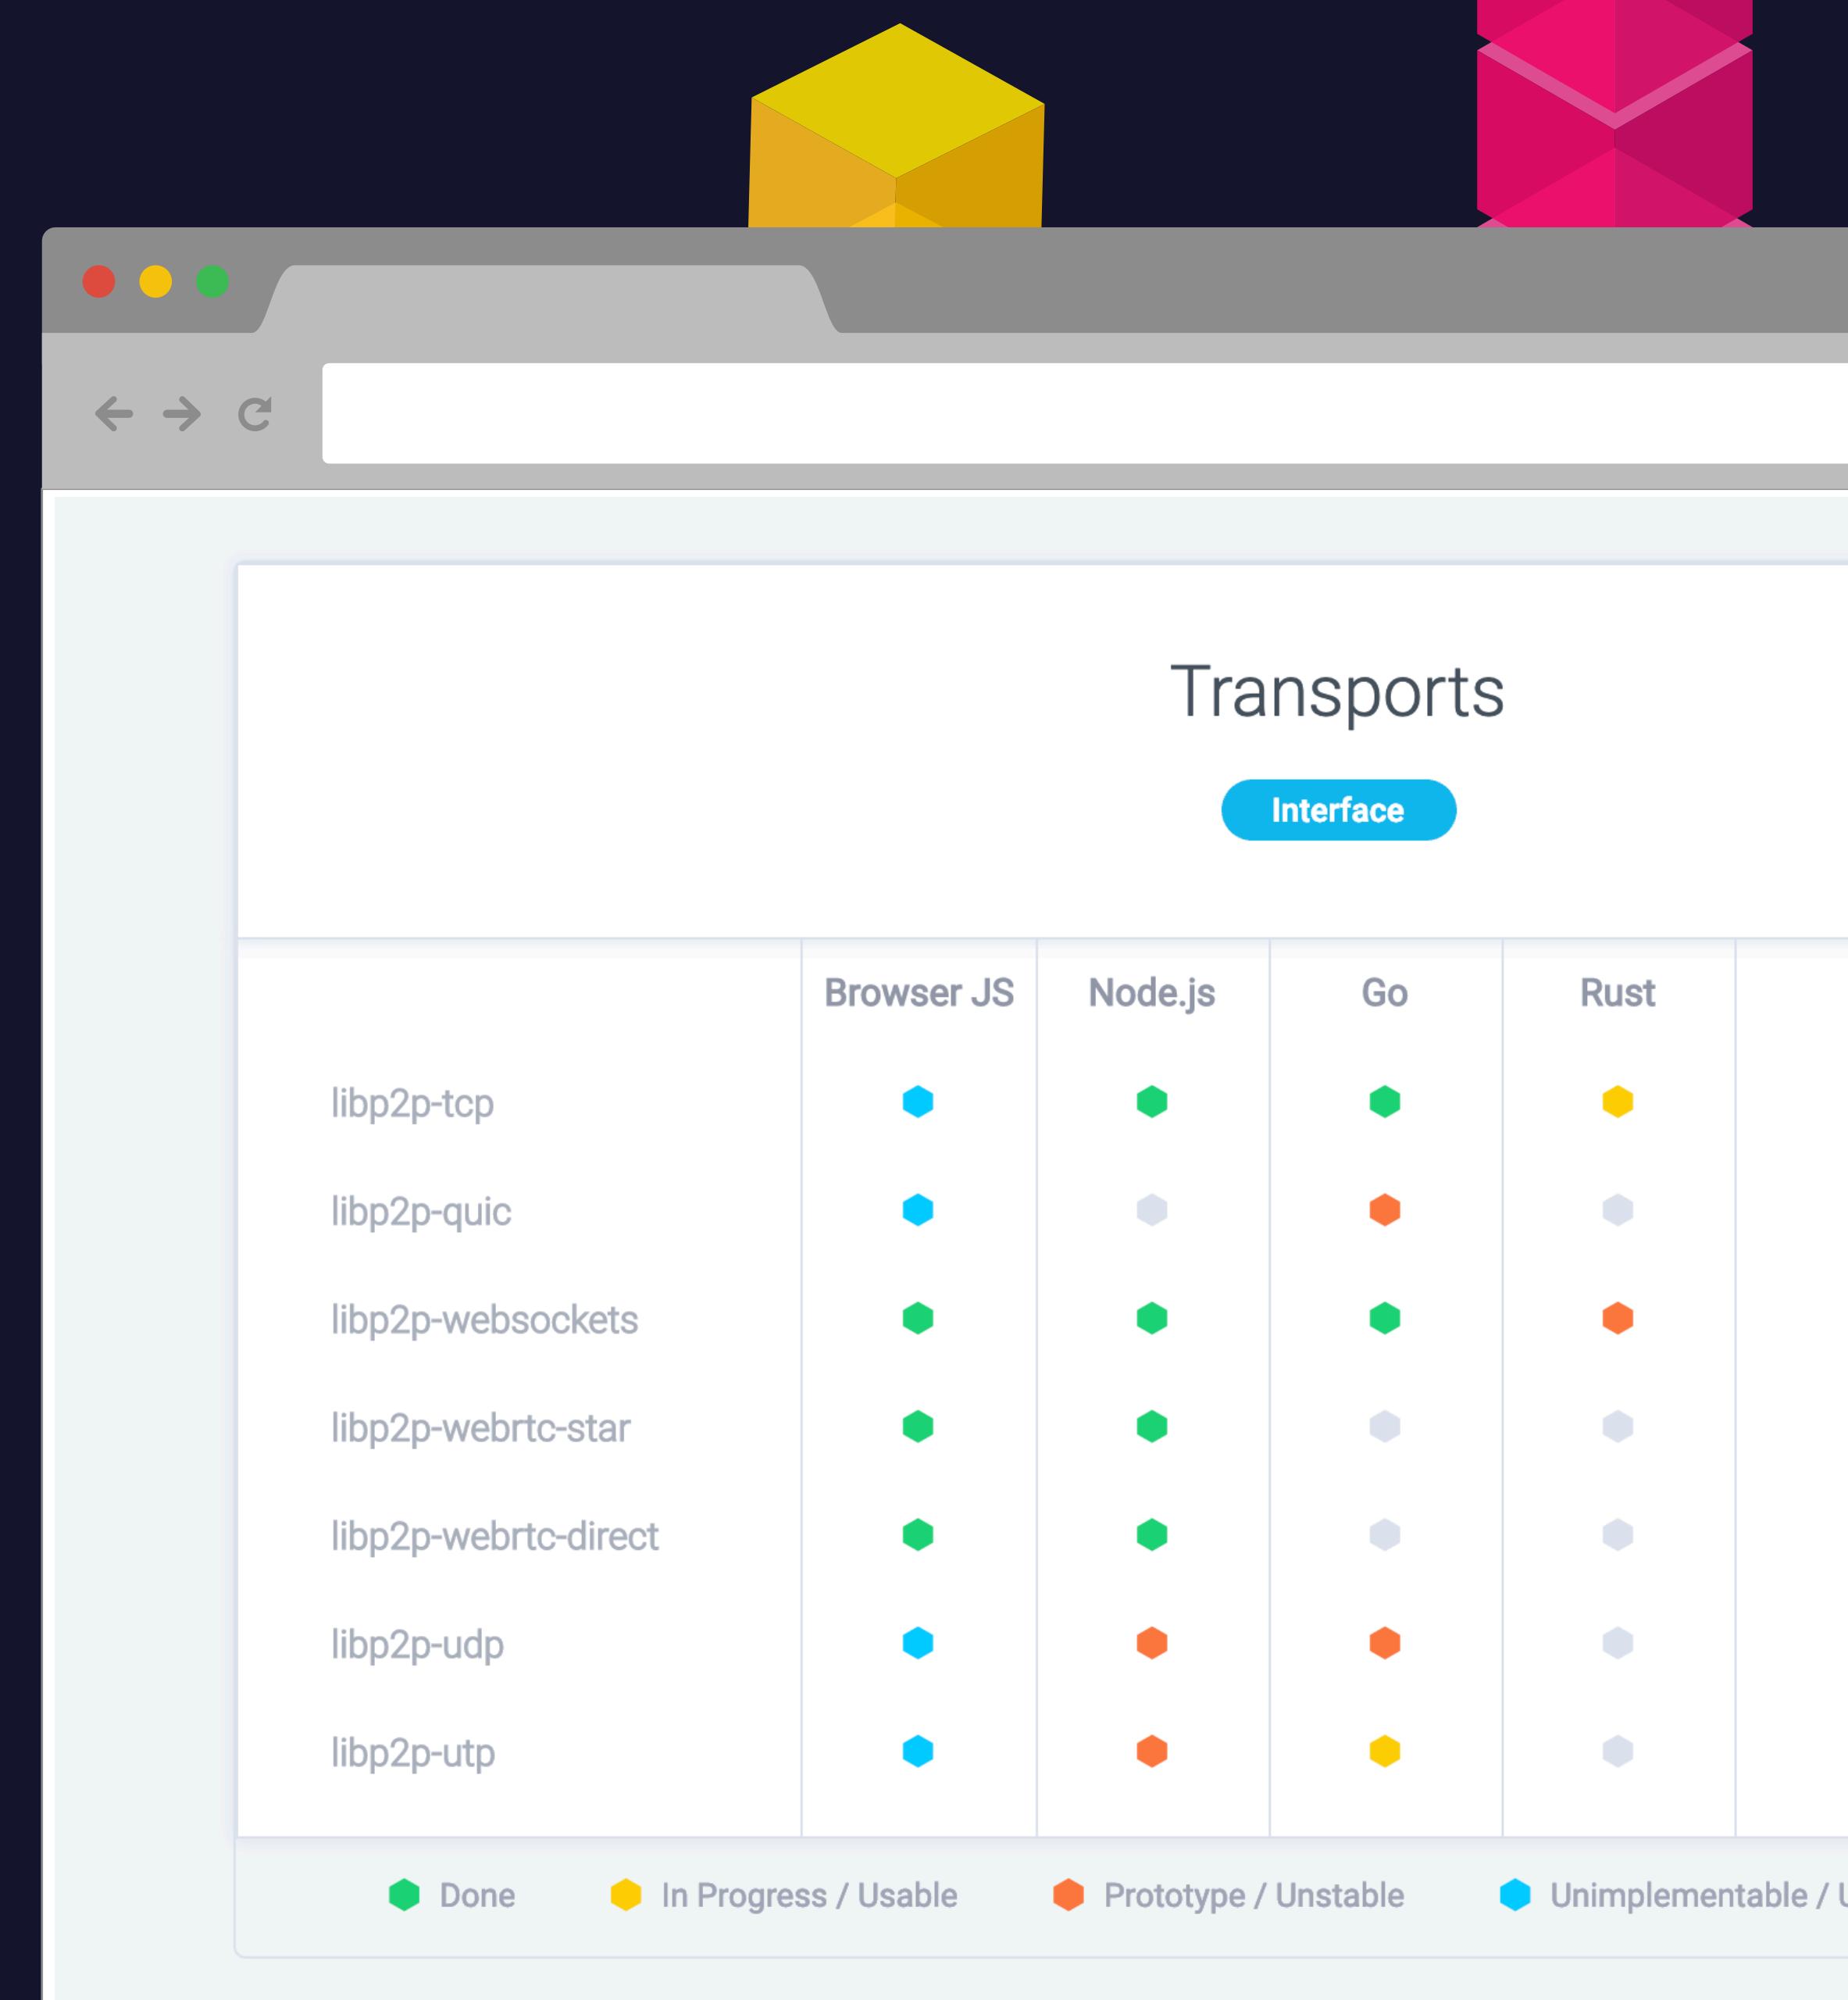




transports

TRANSPORTS

- **transports** are a core abstraction of libp2p.
 - Think of transports as *connection factories*.
 - Dialing and listening.
 - They produce raw connections/sockets/sessions – “the L4” of libp2p.
- Current:
 - TCP
 - QUIC
 - WebSockets
 - WebRTC

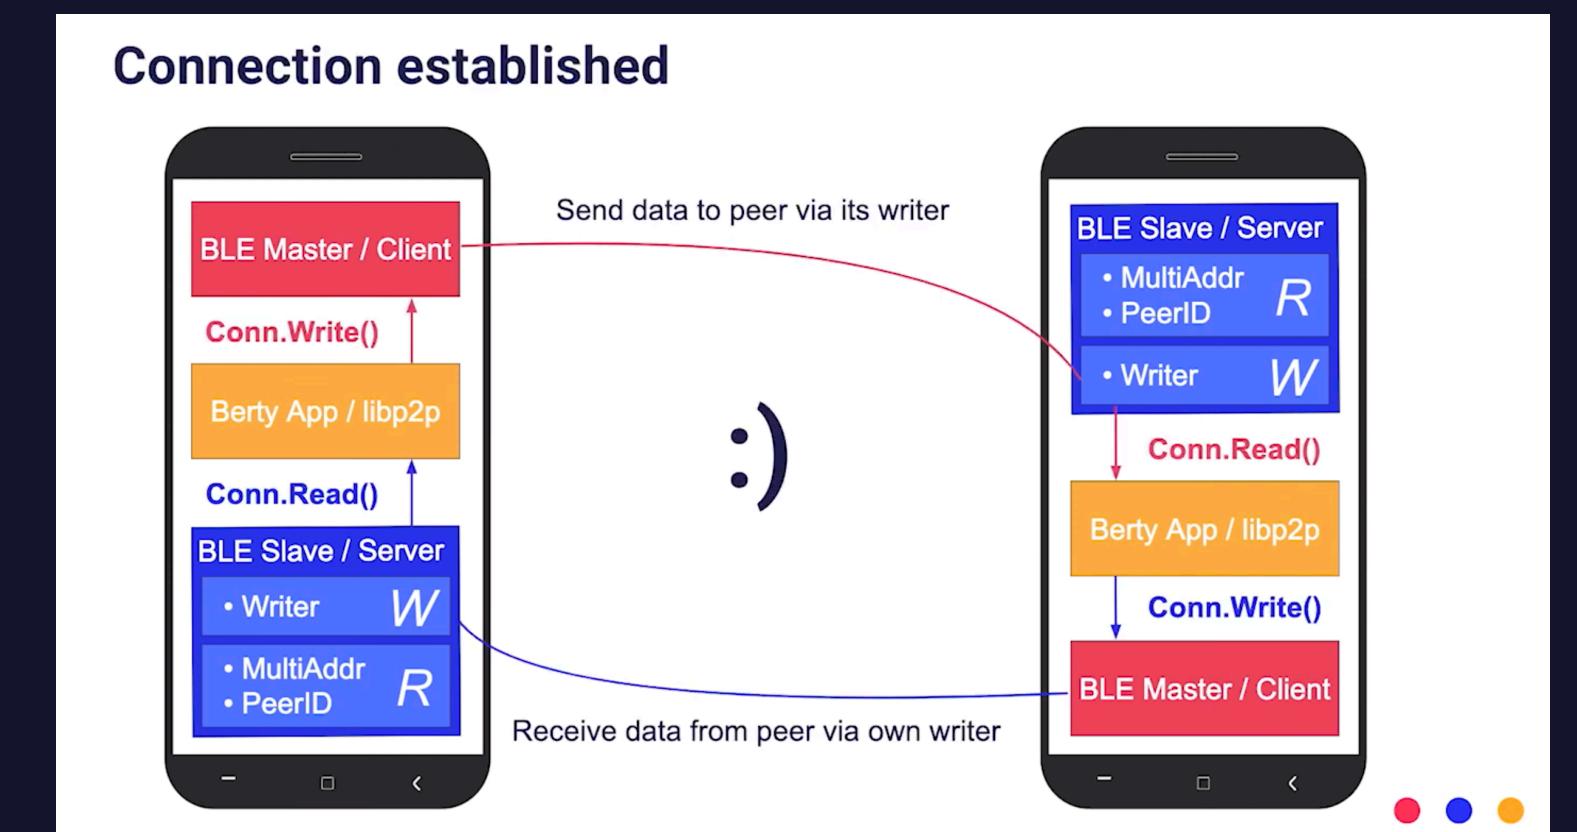
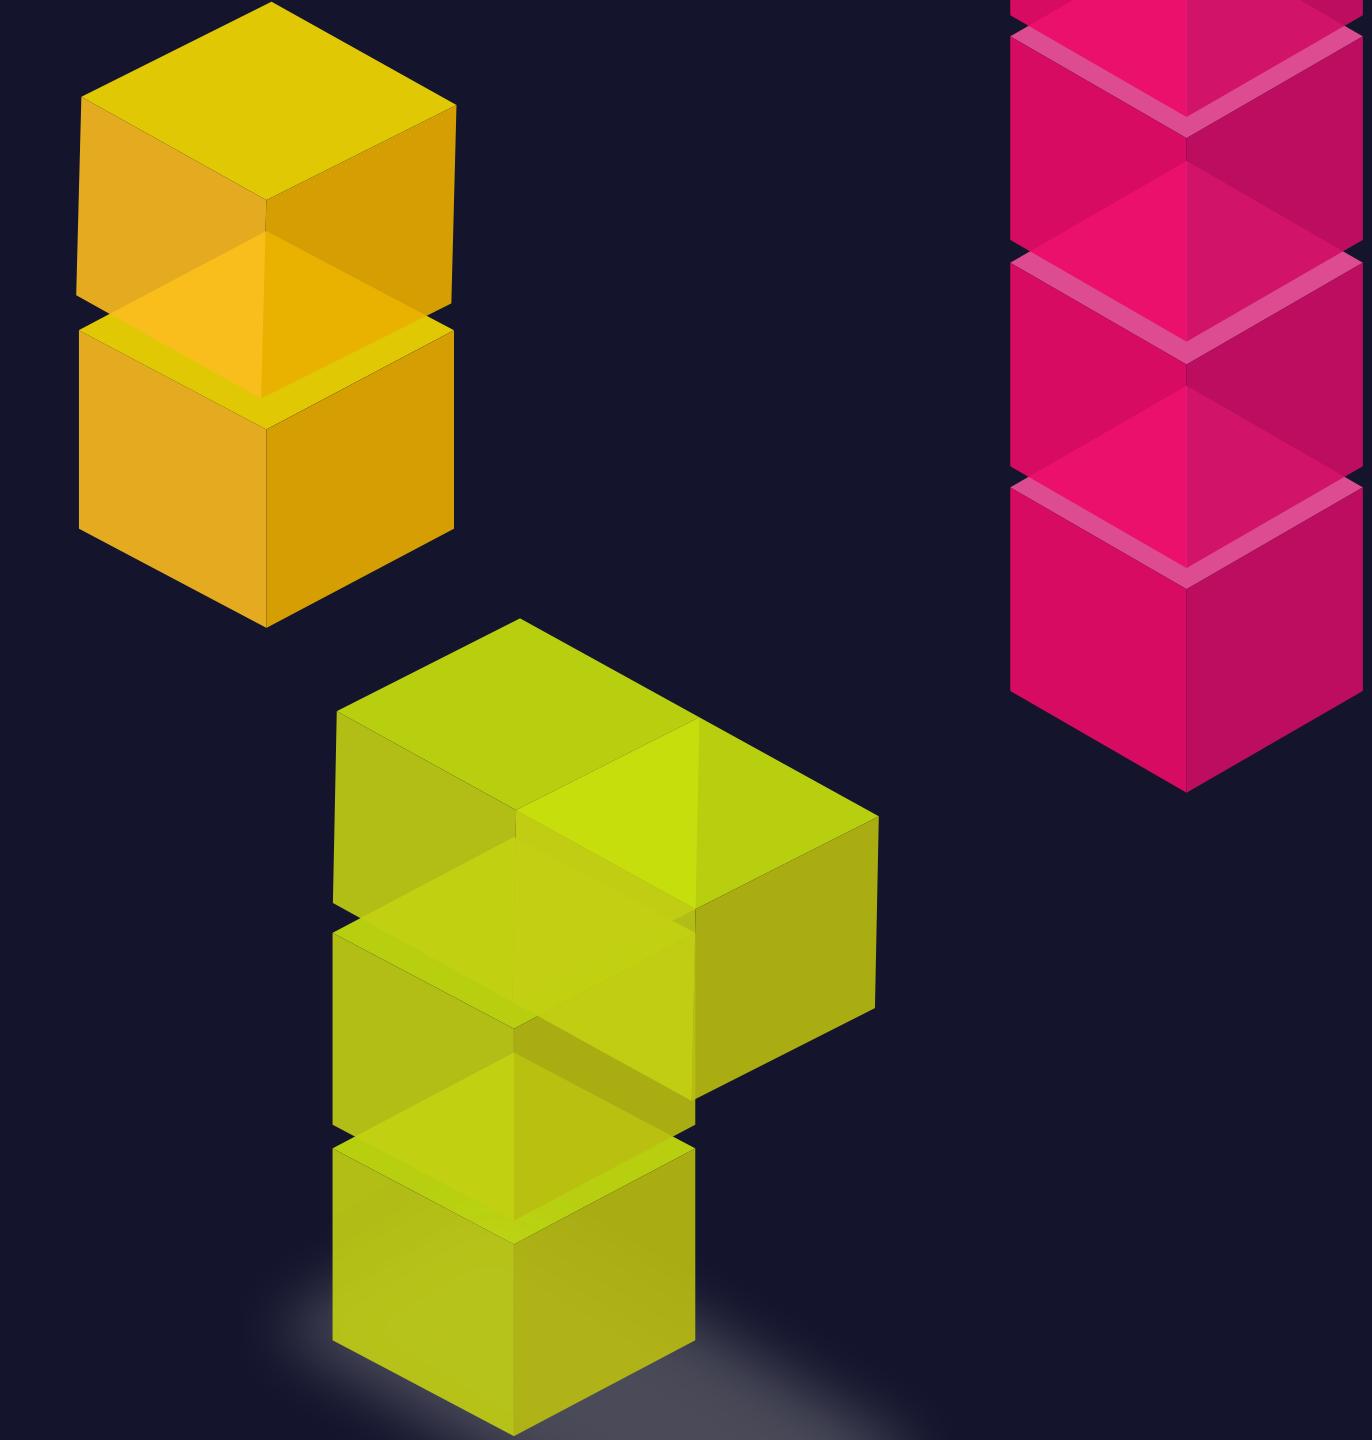


transports

coming soon, to a node near you...

TRANSPORTS

- Stable, ubiquitous QUIC.
- Bluetooth (go, via our friends @ Berty).
- μTP.
- Message orientation: UDP.
- new transports => new multiaddrs
- graceful upgrade path & backwards compatibility.





transports

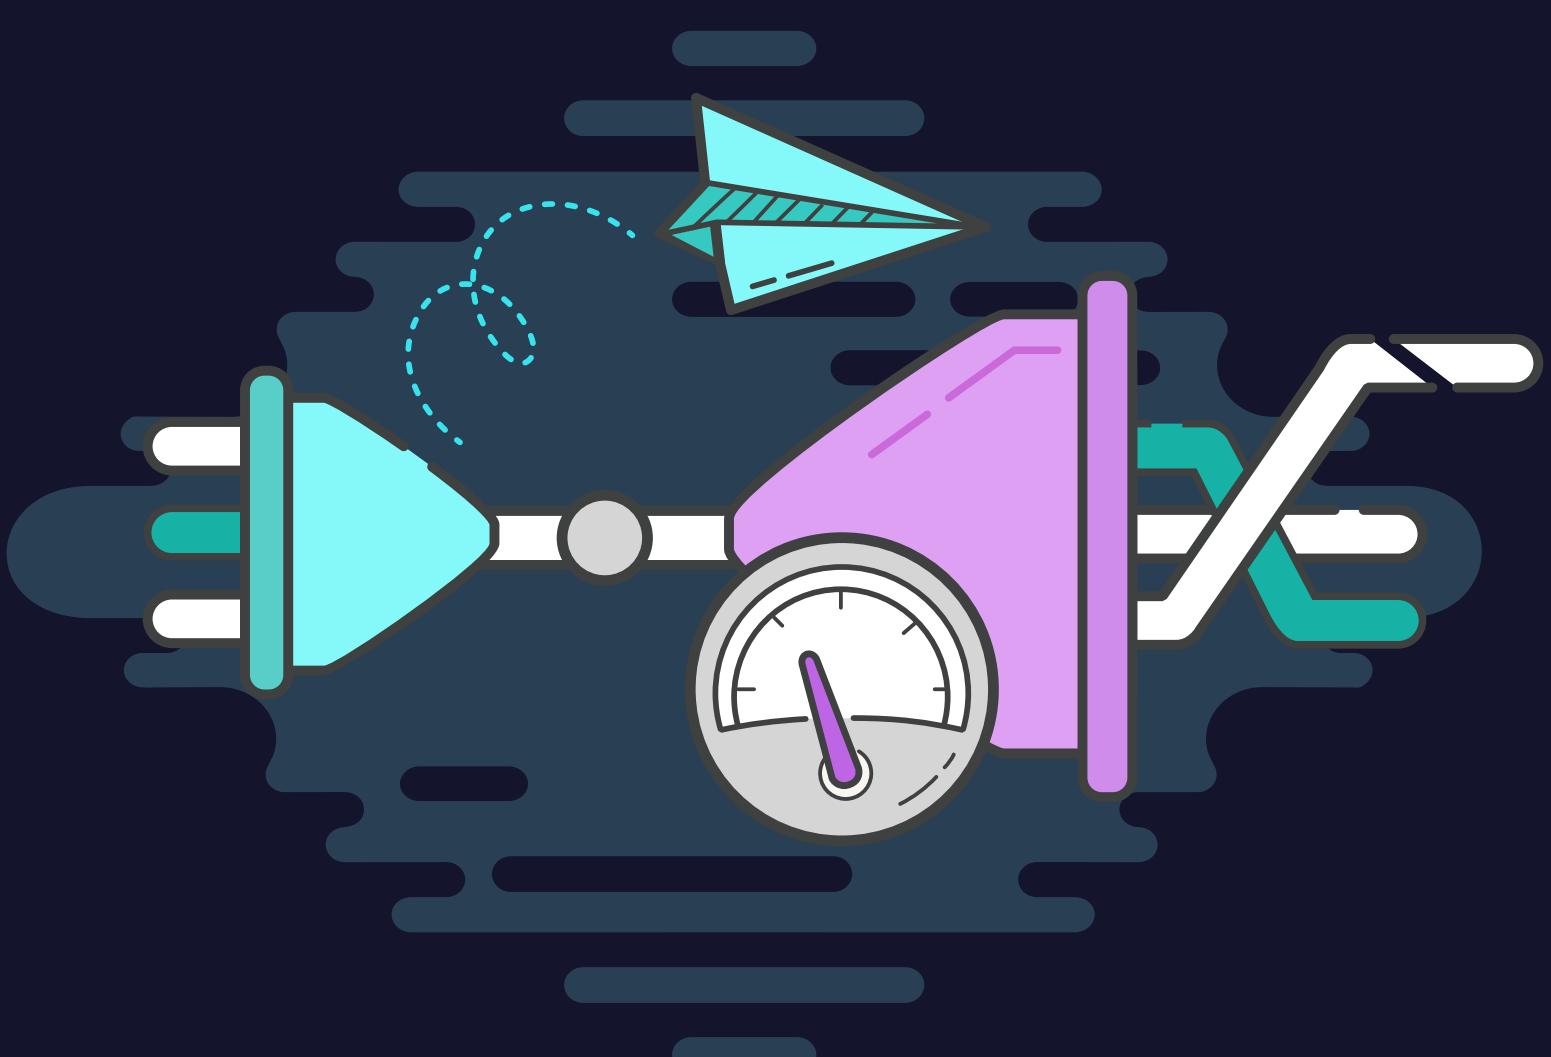
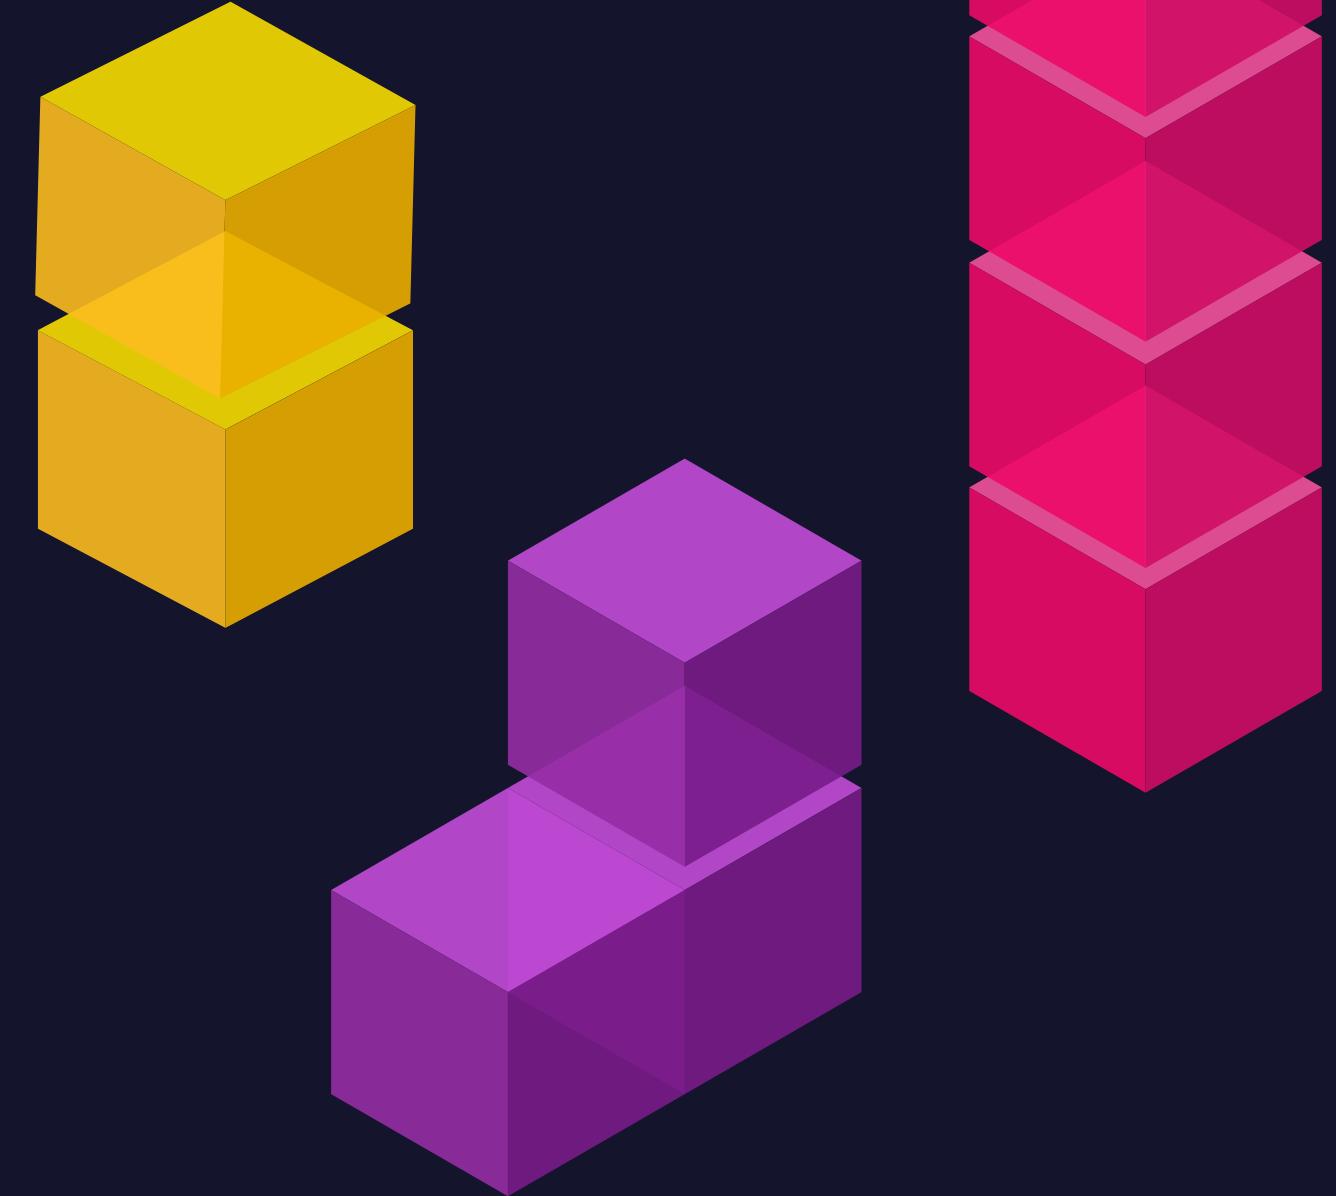
DIALLING...

/ip4/1.2.3.4/tcp/4545/ws/id/QmSoLV6aBf...



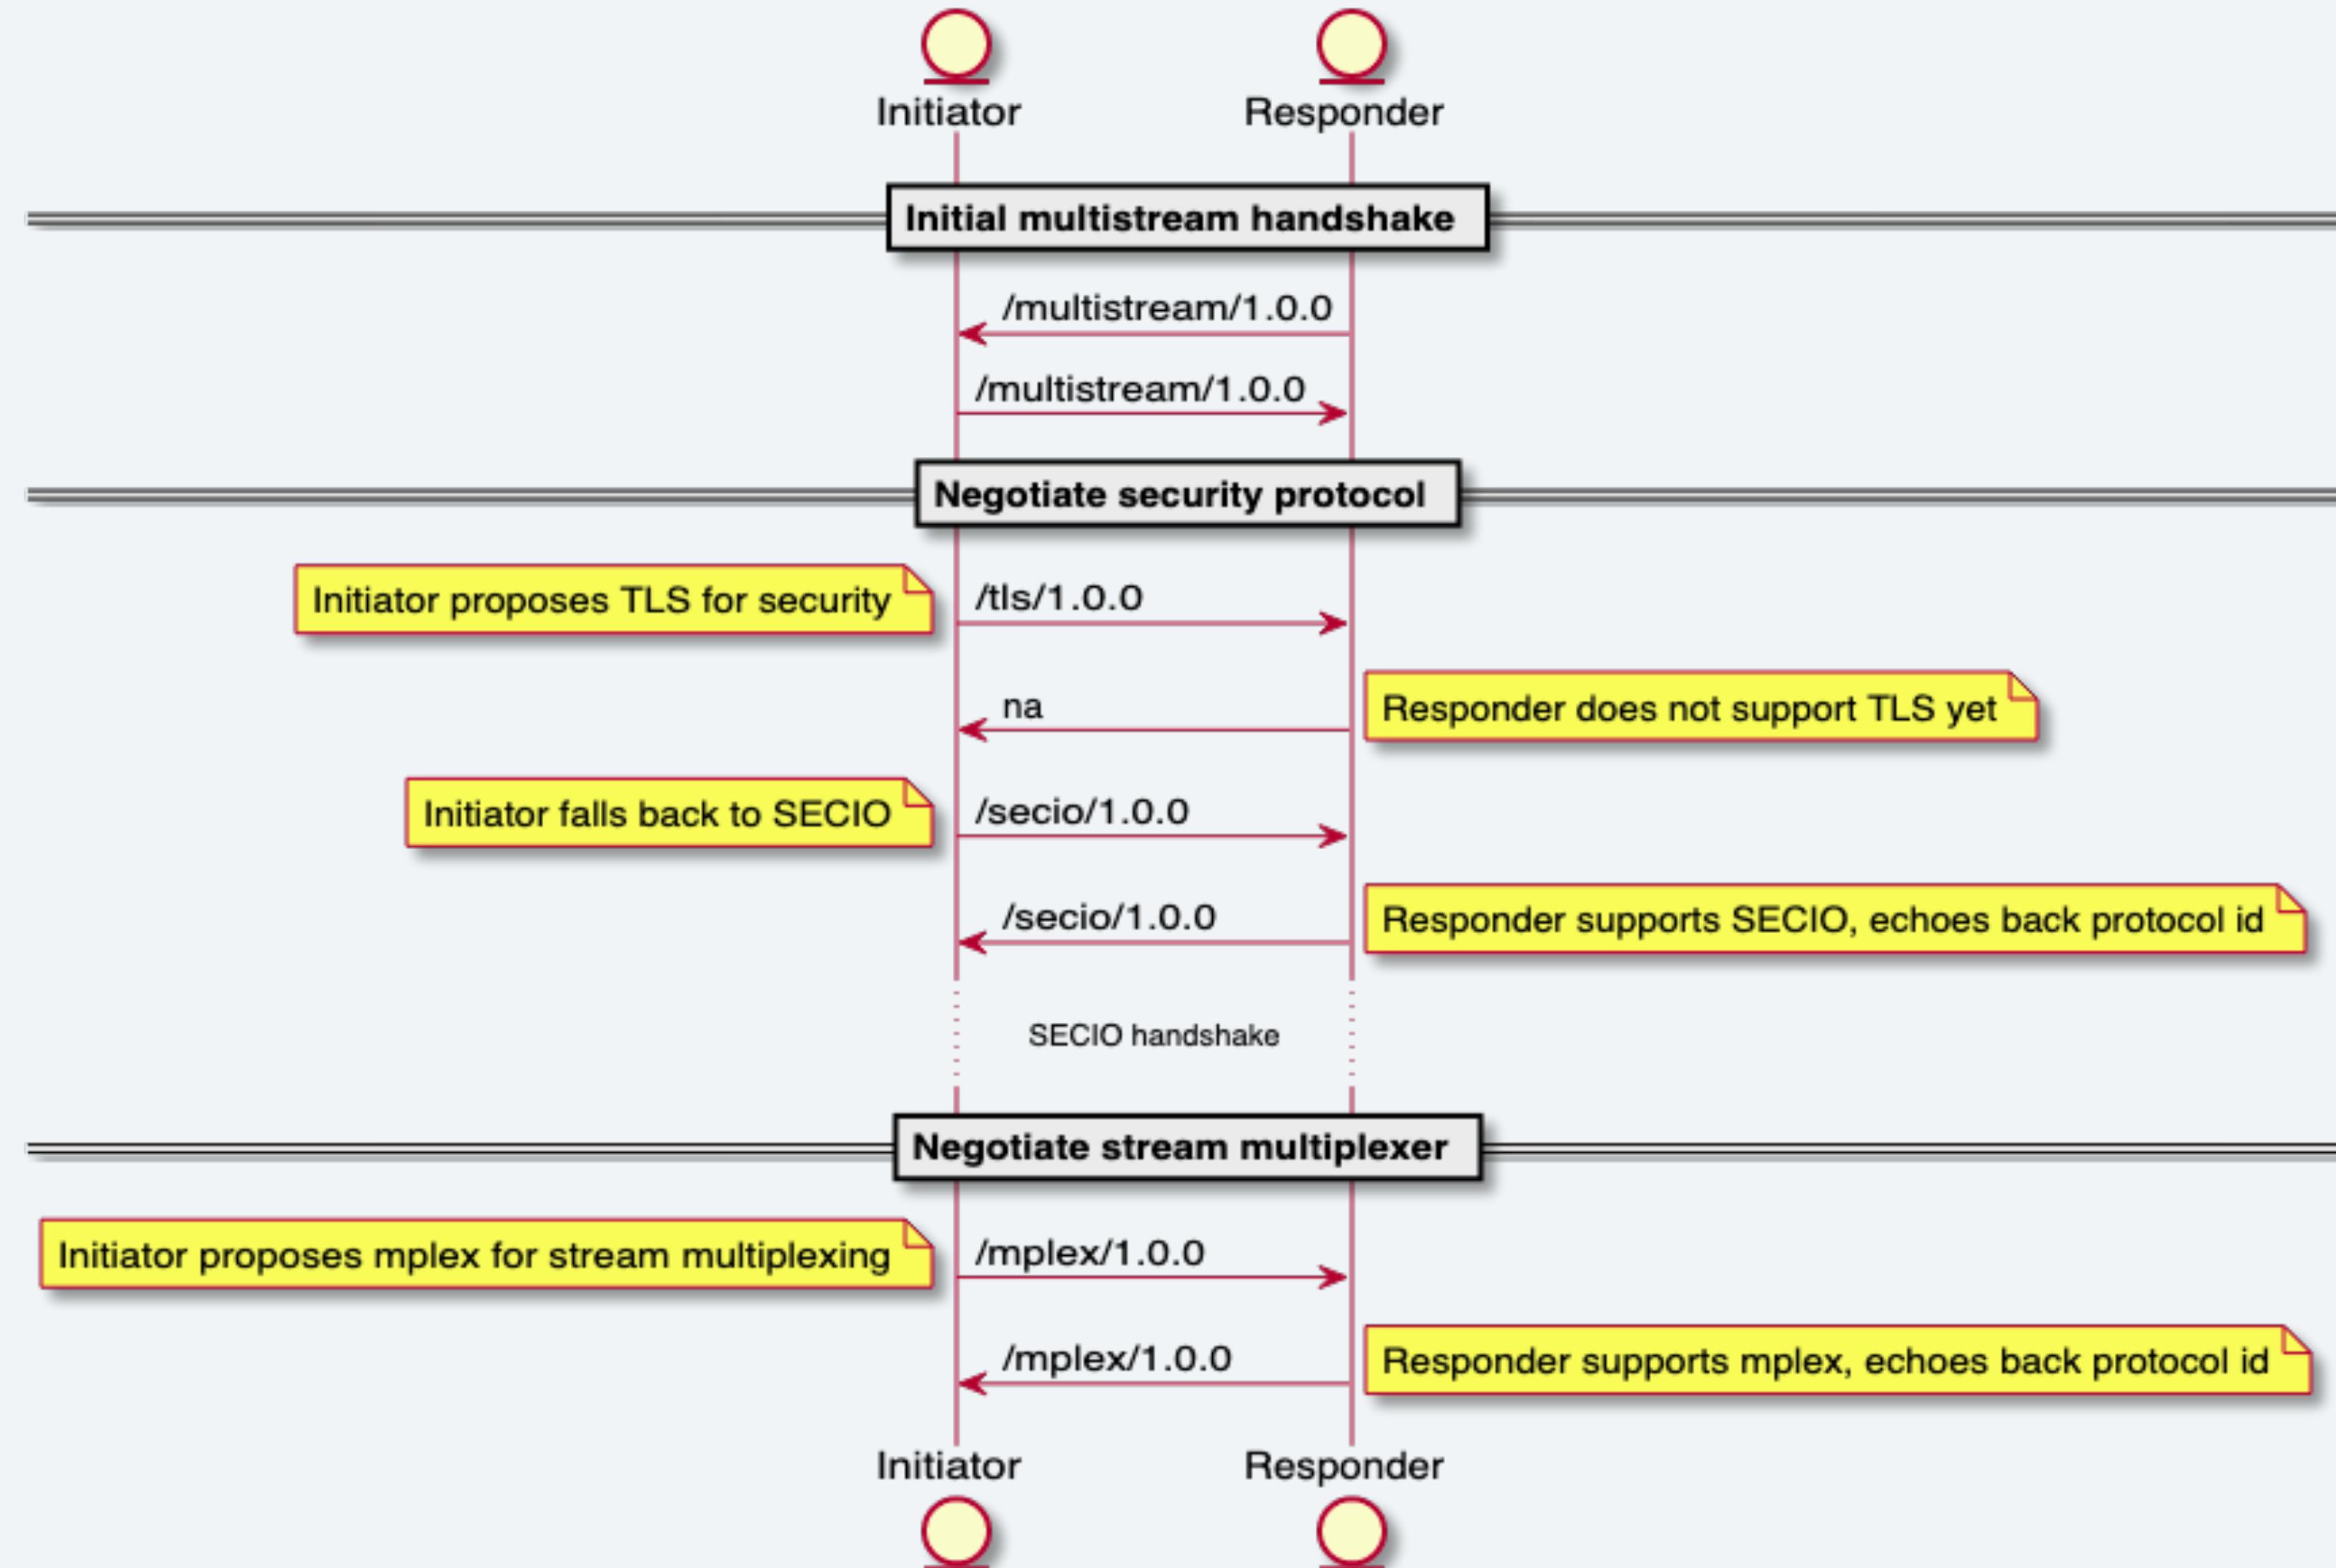
CONNECTION BOOTSTRAPPING

- Connection bootstrapping takes a raw connection, created by a transport, and “*upgrades*” it to a...
- ***Capable connection:*** authenticated, encrypted and multiplexed.
- **We negotiate the secure channel and the multiplexer dynamically.**
 - multistream-select 1.0; more on this later.
 - Some transports don’t need *upgrading*.
 - They’re natively capable of encryption and multiplexing.
 - QUIC is an example.





connection bootstrapping



<https://github.com/libp2p/specs/blob/master/connections/README.md>



secure channels

SECURE CHANNELS

- Peer authentication.
- Transport encryption.
 - Does not preclude you from encrypting/signing application data.
- Transmission integrity. Non-repudiation. HMAC.
- All implementations: SECIO.
 - Some: TLS 1.3.
- For testing:
 - /plaintext/2.0.0. Try it out with Wireshark! 🎈



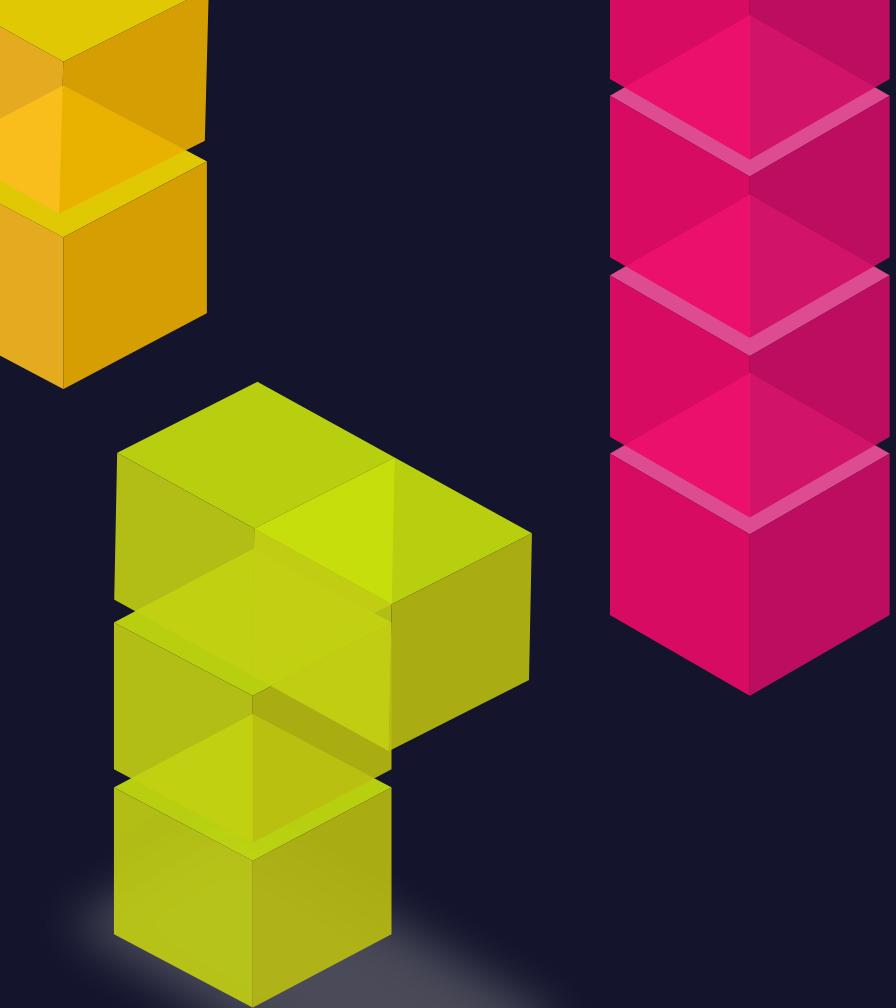


secure channels

coming soon, to a node near you...

SECURE CHANNELS

- new secure channels: **TLS 1.3 & Noise handshakes.**
- **Noise Handshakes.**
 - Noise Pipes: IK for optimistic scenario, falling back to XX.
 - Encrypted early data to expedite multiplexer negotiation.
 - Required by Ethereum 2.0 mainnet.
- **TLS 1.3 in all languages.**
 - Available in go-libp2p. Can enable as an experiment in IPFS.
 - Roadblock in rust-libp2p, **devgrant** available for js-libp2p!



xx:
-> e
<- e, ee, s, es
-> s, se

IK:
<- s
...
-> e, es, s, ss
<- e, ee, se

xxfallback:
-> e
...
<- e, ee, s, es
-> s, se



multiplexing

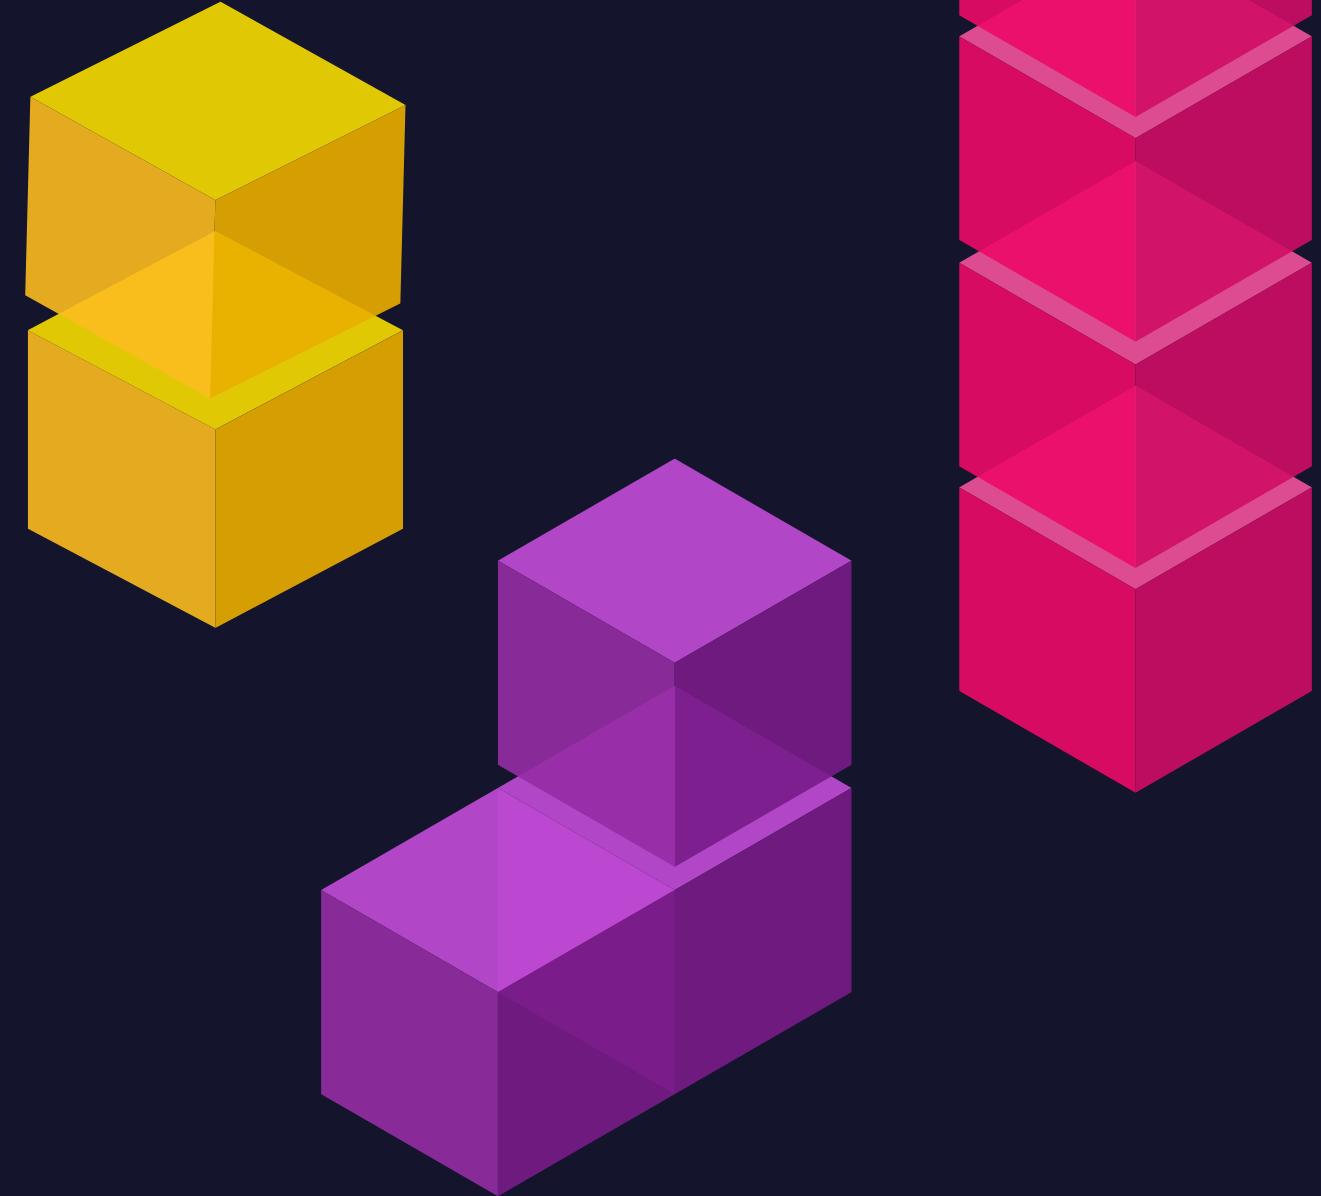
holding many conversations at once

MULTIPLEXING



Establishing a p2p connection may not be cheap or easy (e.g. hole punching, negotiation, handshake, etc.)

Let's amortise it.





multiplexing

holding many conversations at once

MULTIPLEXING

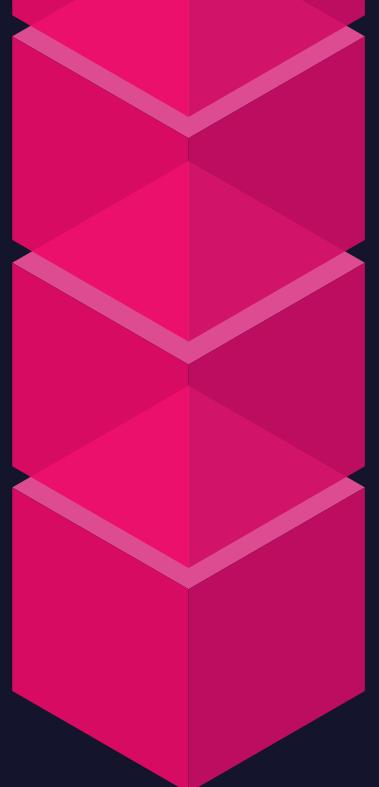
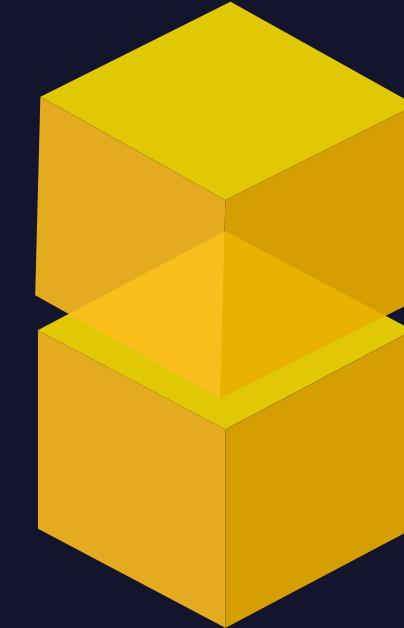


- yamux (hashicorp)
- mplex (libp2p)
- spdystream
- quic native multiplexing

streams are identified by a number.
how do we contextualise the stream to a particular protocol?



future: monplexing





protocol negotiation

multistream-select 1.0

PROTOCOL NEGOTIATION

- multistream-select 1.0 is our current protocol negotiation scheme.
- interactive, chatty, speculative, and wasteful in pessimistic scenarios.

```
>>> /multistream/1.0.0          (1)
<<< /multistream/1.0.0          (2)
>>> /foo/A                      (3)
<<< na                         (4)
>>> /foo/B                      (5)
<<< /foo/B                      (6)
```





protocol negotiation

towards multiselect 2.0

PROTOCOL NEGOTIATION

- **multiselect 2.0 for efficient protocol negotiation.**
 - <https://github.com/libp2p/specs/pull/205>, features:
 - 1-RTT connection bootstrapping.
 - cheap repetitive protocol selection.
 - TCP simultaneous connect support.
 - session resumption.
 - upfront protocol tables.
 - confident selection. interoperable with deterministic or probabilistic discovery-level advertisements, e.g. bloom filters.

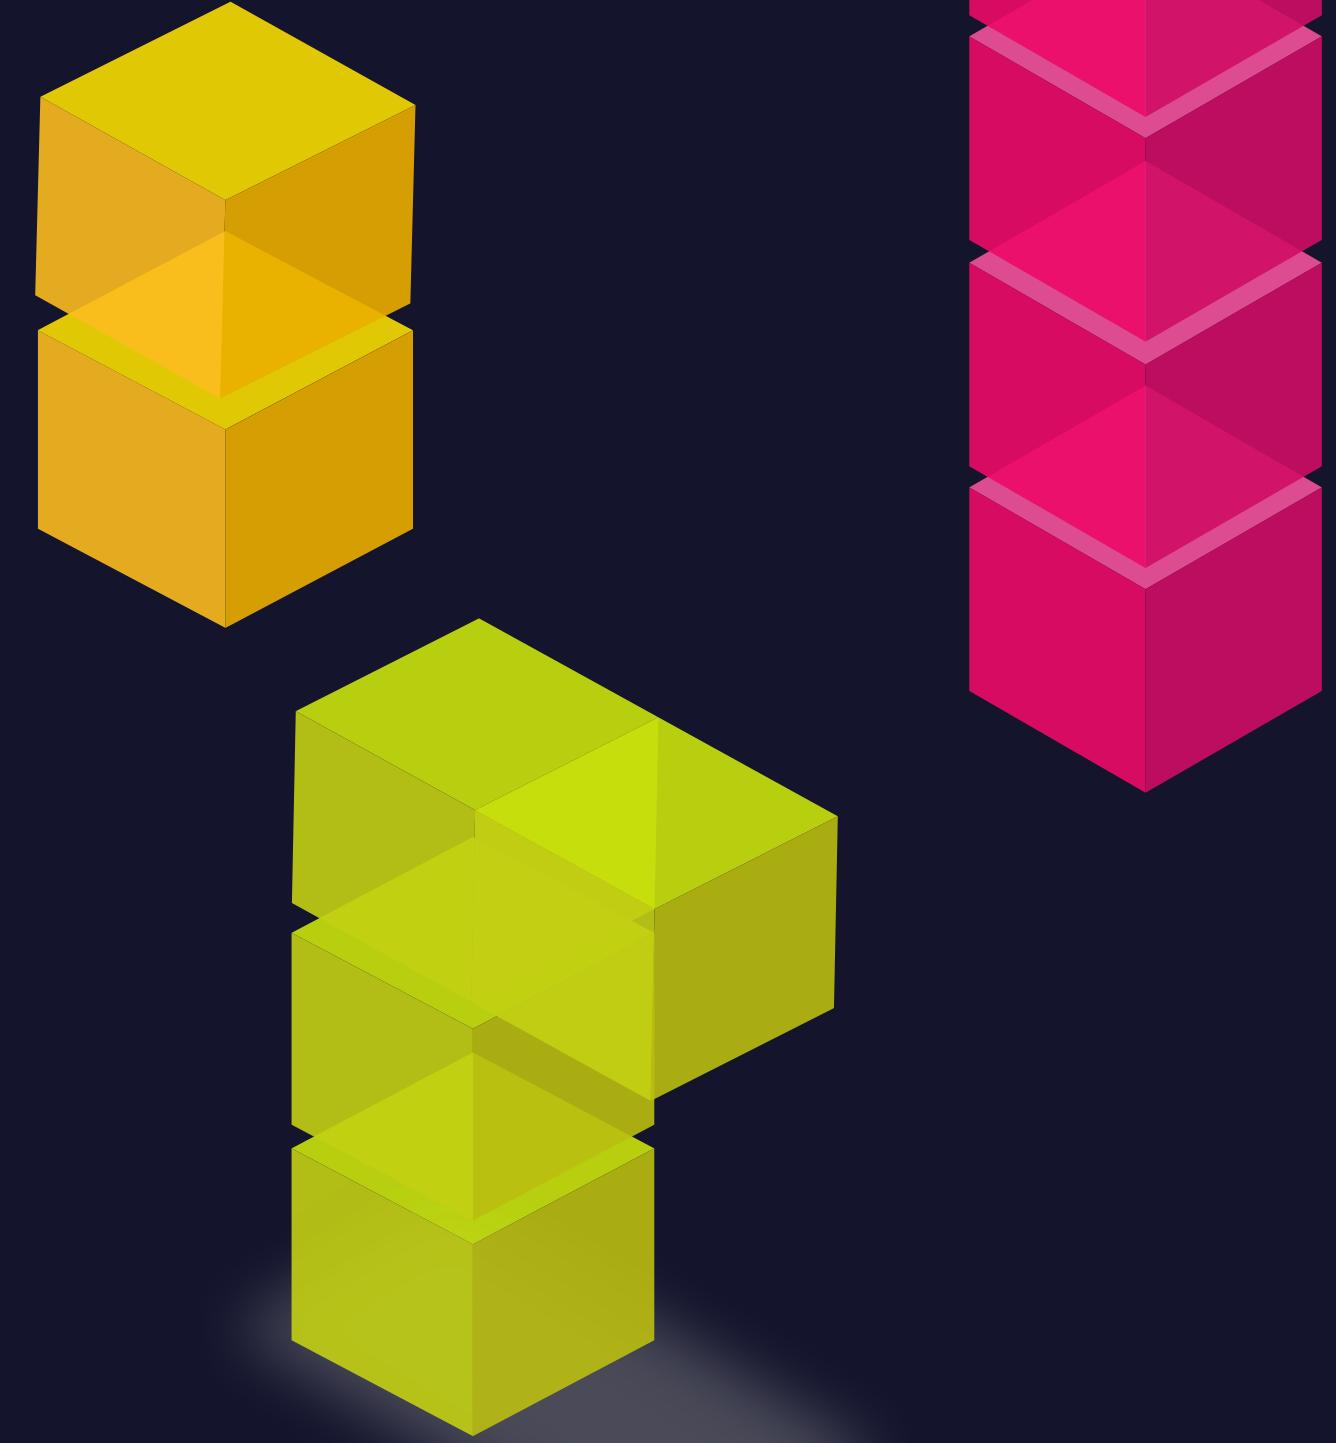




protocol negotiation

PROTOCOL HANDLING

- **Protocol IDs:**
 - `/floodsub/1.0.0`
 - `/ipfs/bitswap/1.0.0`
 - `/eth2/beacon_chain/req/...`
- analogous to REST endpoints, they serve for routing and handling.
 - can represent an entire protocol, or a message type.
 - to attach a new protocol, you mount a handler.
- semver under question – we'll probably migrate away.





PEER DISCOVERY



```
// Advertiser is an interface for advertising services
type Advertiser interface {
    // Advertise advertises a service
    Advertise(ctx context.Context, ns string, opts ...Option) (time.Duration, error)
}

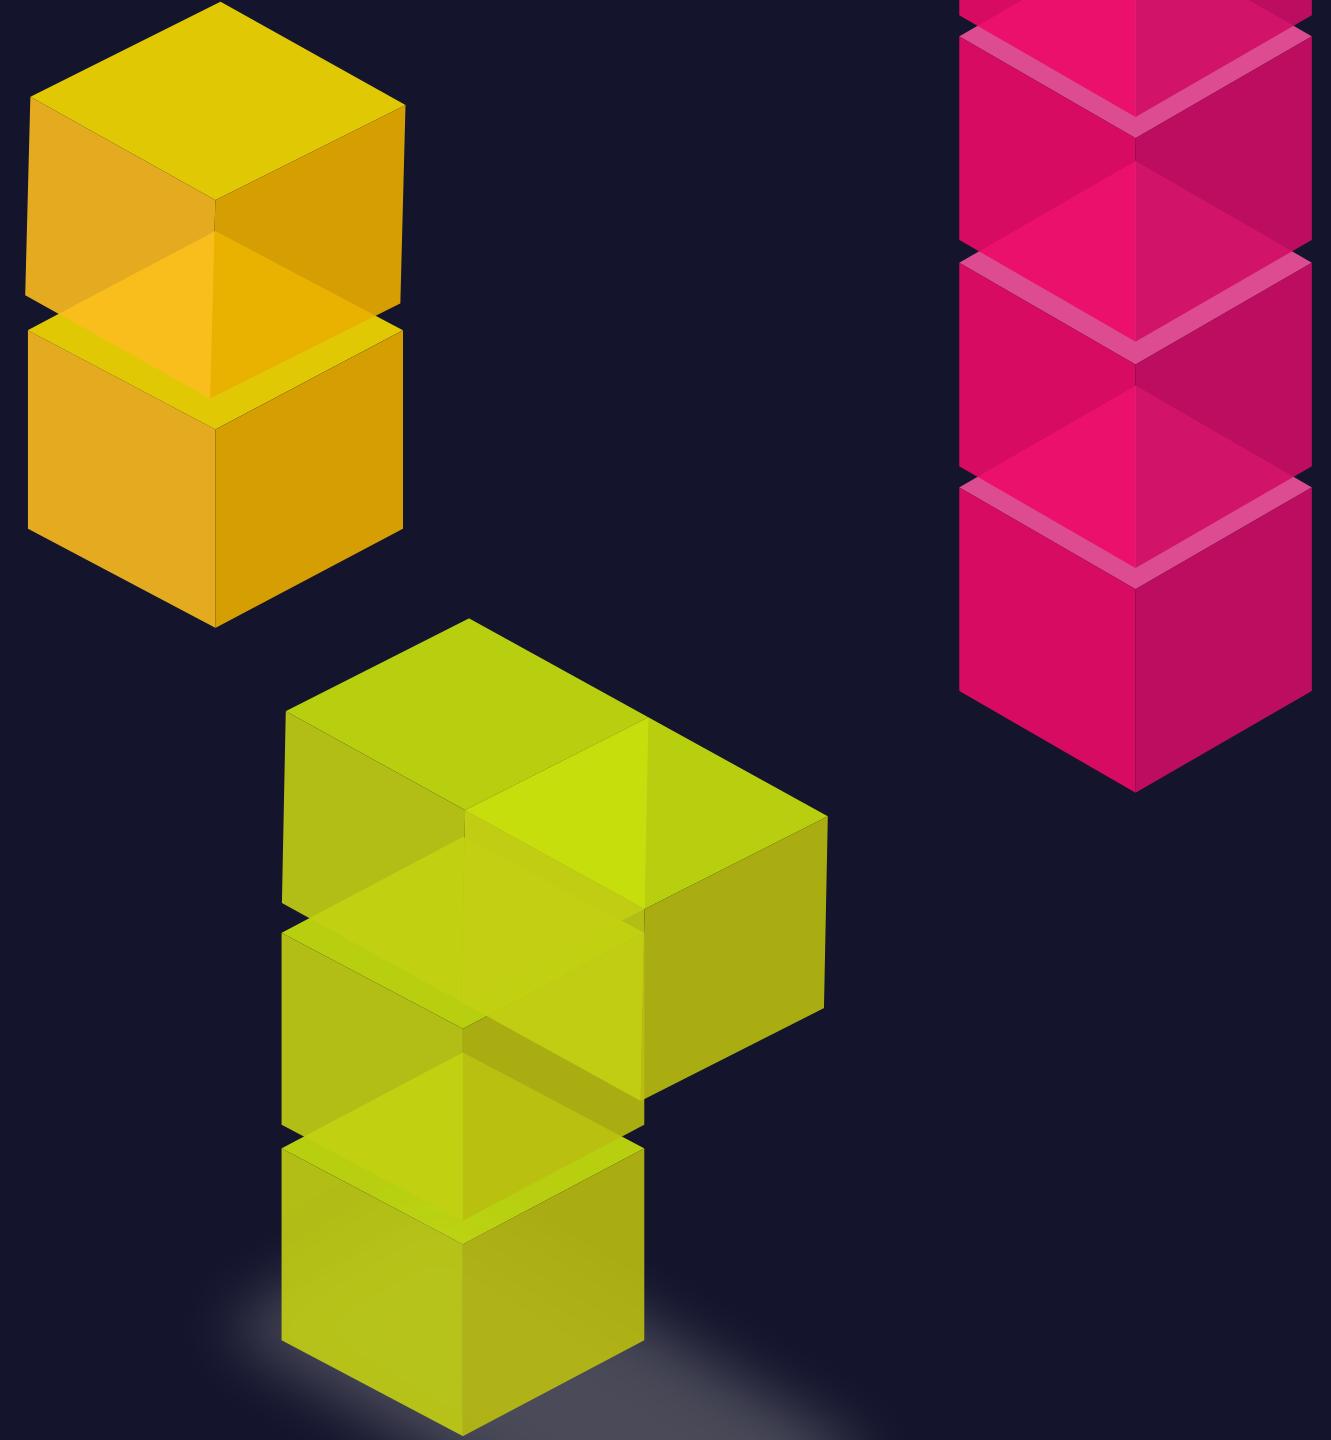
// Discoverer is an interface for peer discovery
type Discoverer interface {
    // FindPeers discovers peers providing a service
    FindPeers(ctx context.Context, ns string, opts ...Option) (<-chan peer.AddrInfo, error)
}

// Discovery is an interface that combines service advertisement and peer discovery
type Discovery interface {
    Advertiser
    Discoverer
}
```



PEER DISCOVERY

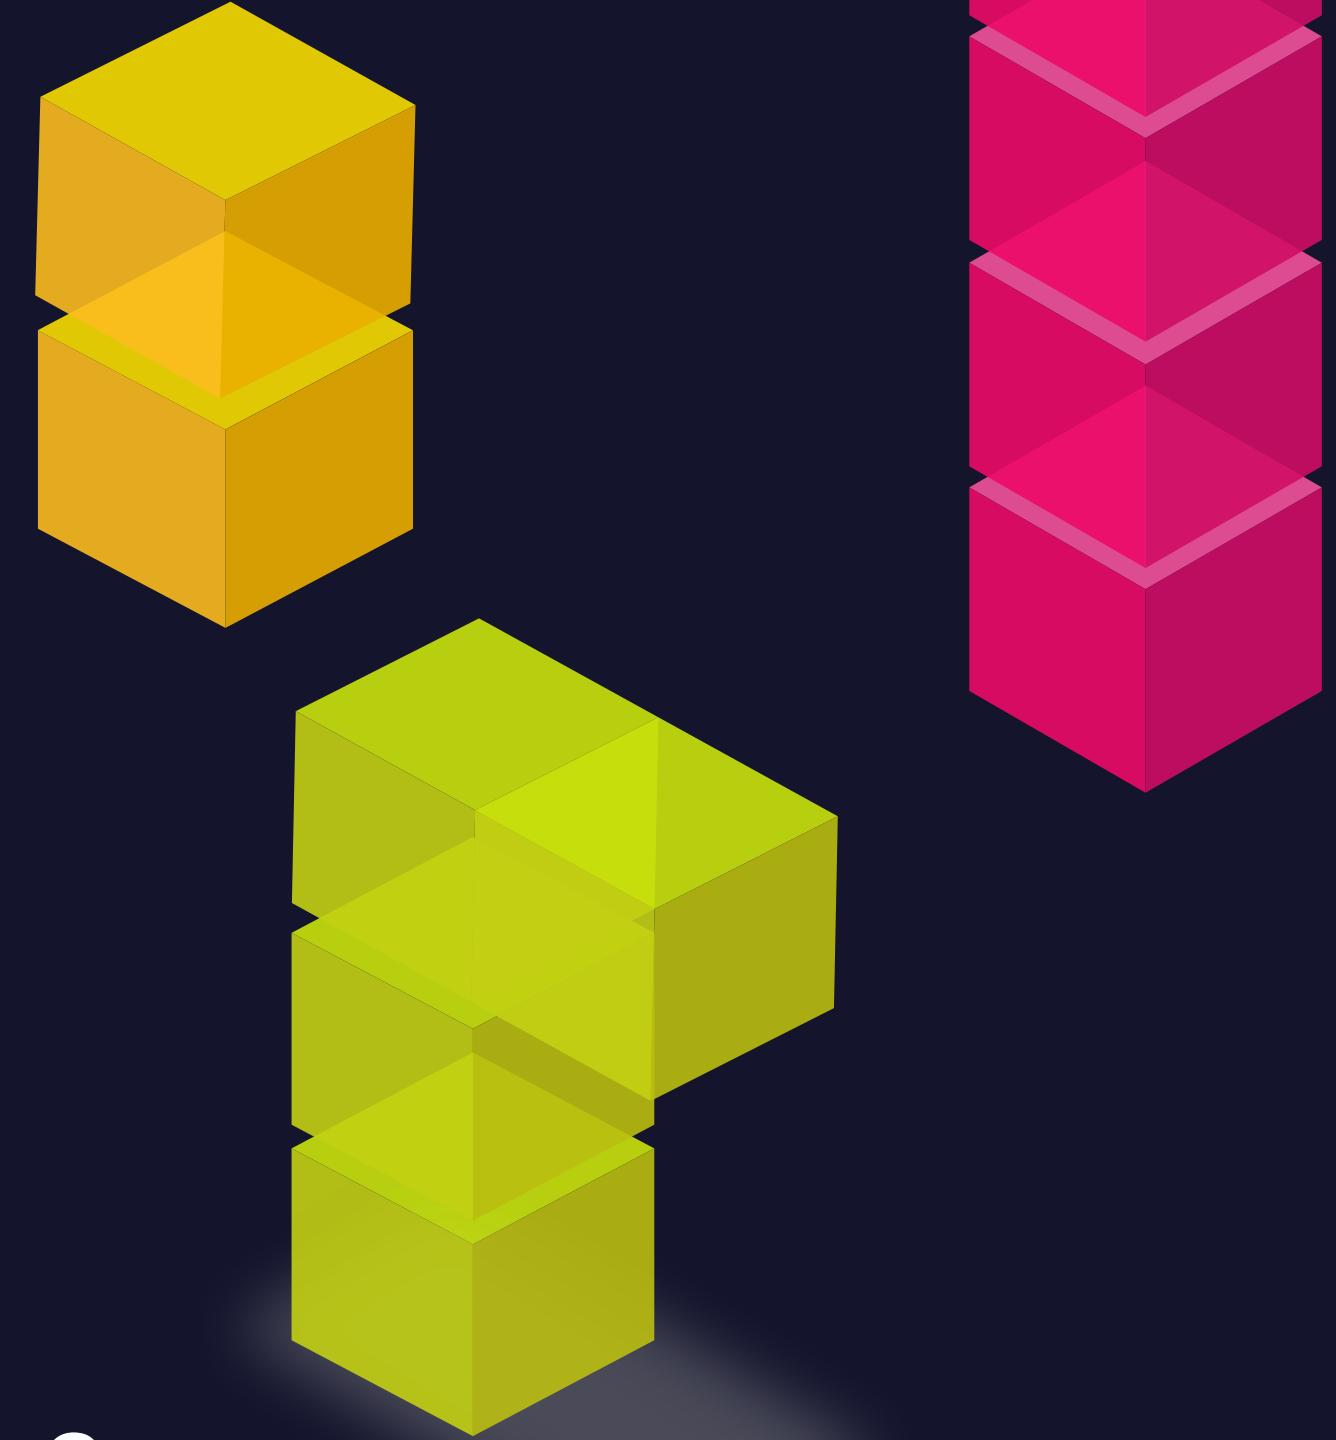
- **backends:**
 - DHT.
 - Rendezvous nodes / signalling servers.
 - mDNS for local peer discovery.
- future: Bluetooth, NFC, etc.
- future: Brahms (membership sampling via pubsub).
- anything, really: matrix channel, IRC, etc.
- **the bootstrapping/seeding issue:** hybrid approaches?





CONTENT ROUTING

- Store, find, advertise content-addressed chunks of data in a decentralised p2p network.
- Also an interface, mostly fulfilled by the DHT.
- Alternatives implementations can query DNS, use pubsub, etc.

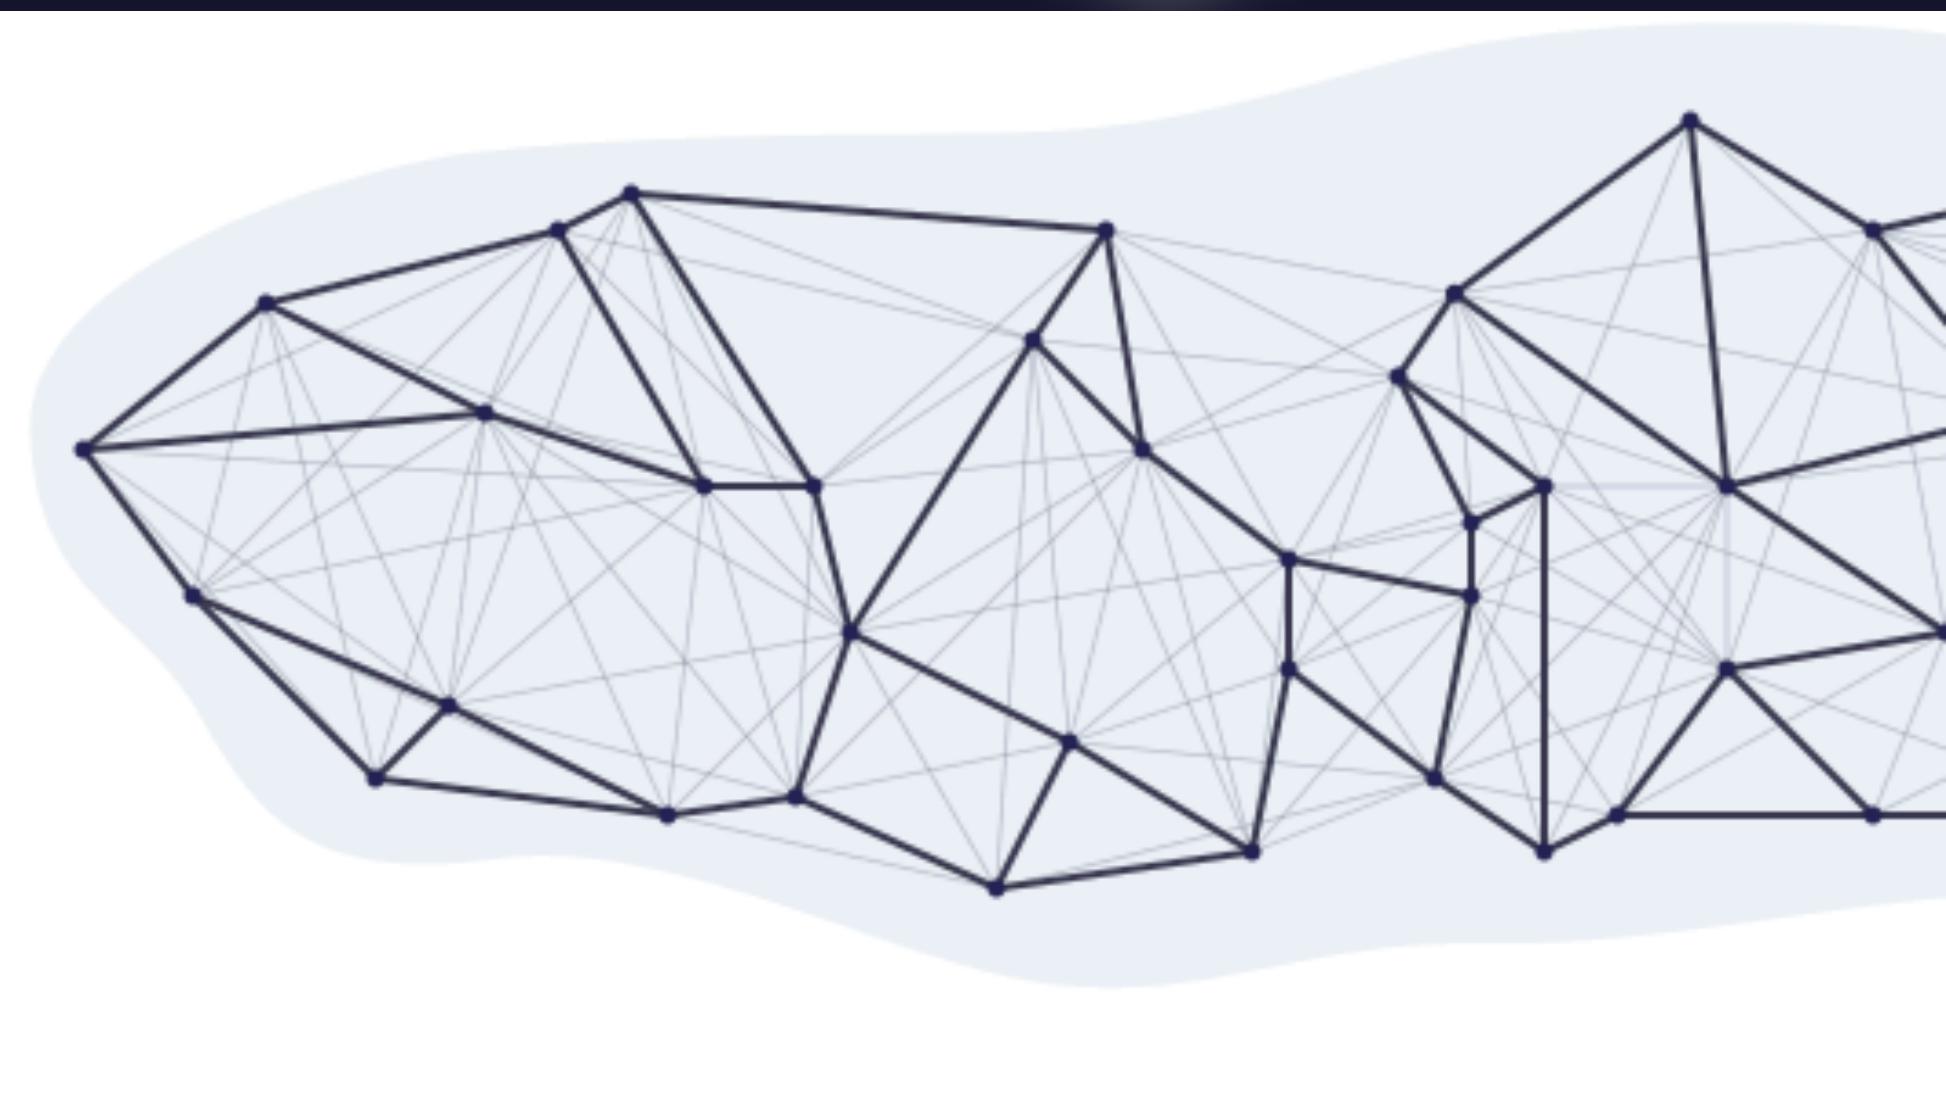
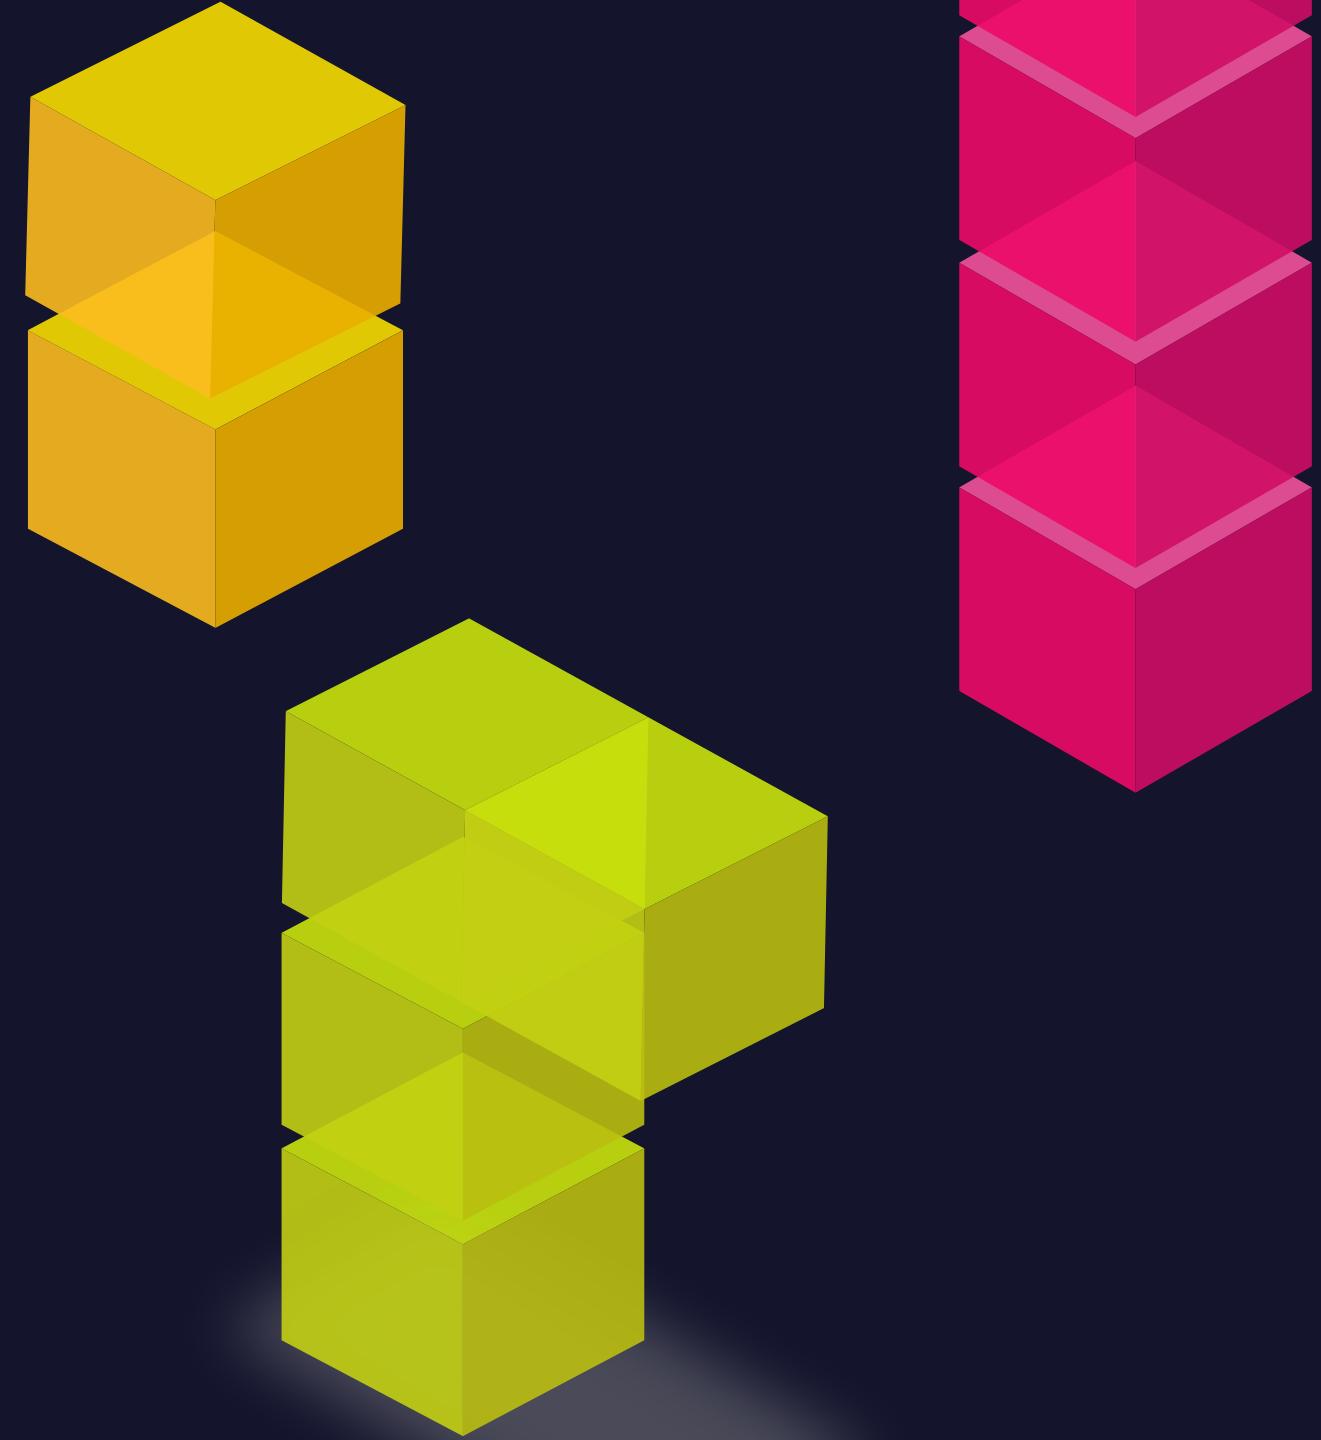




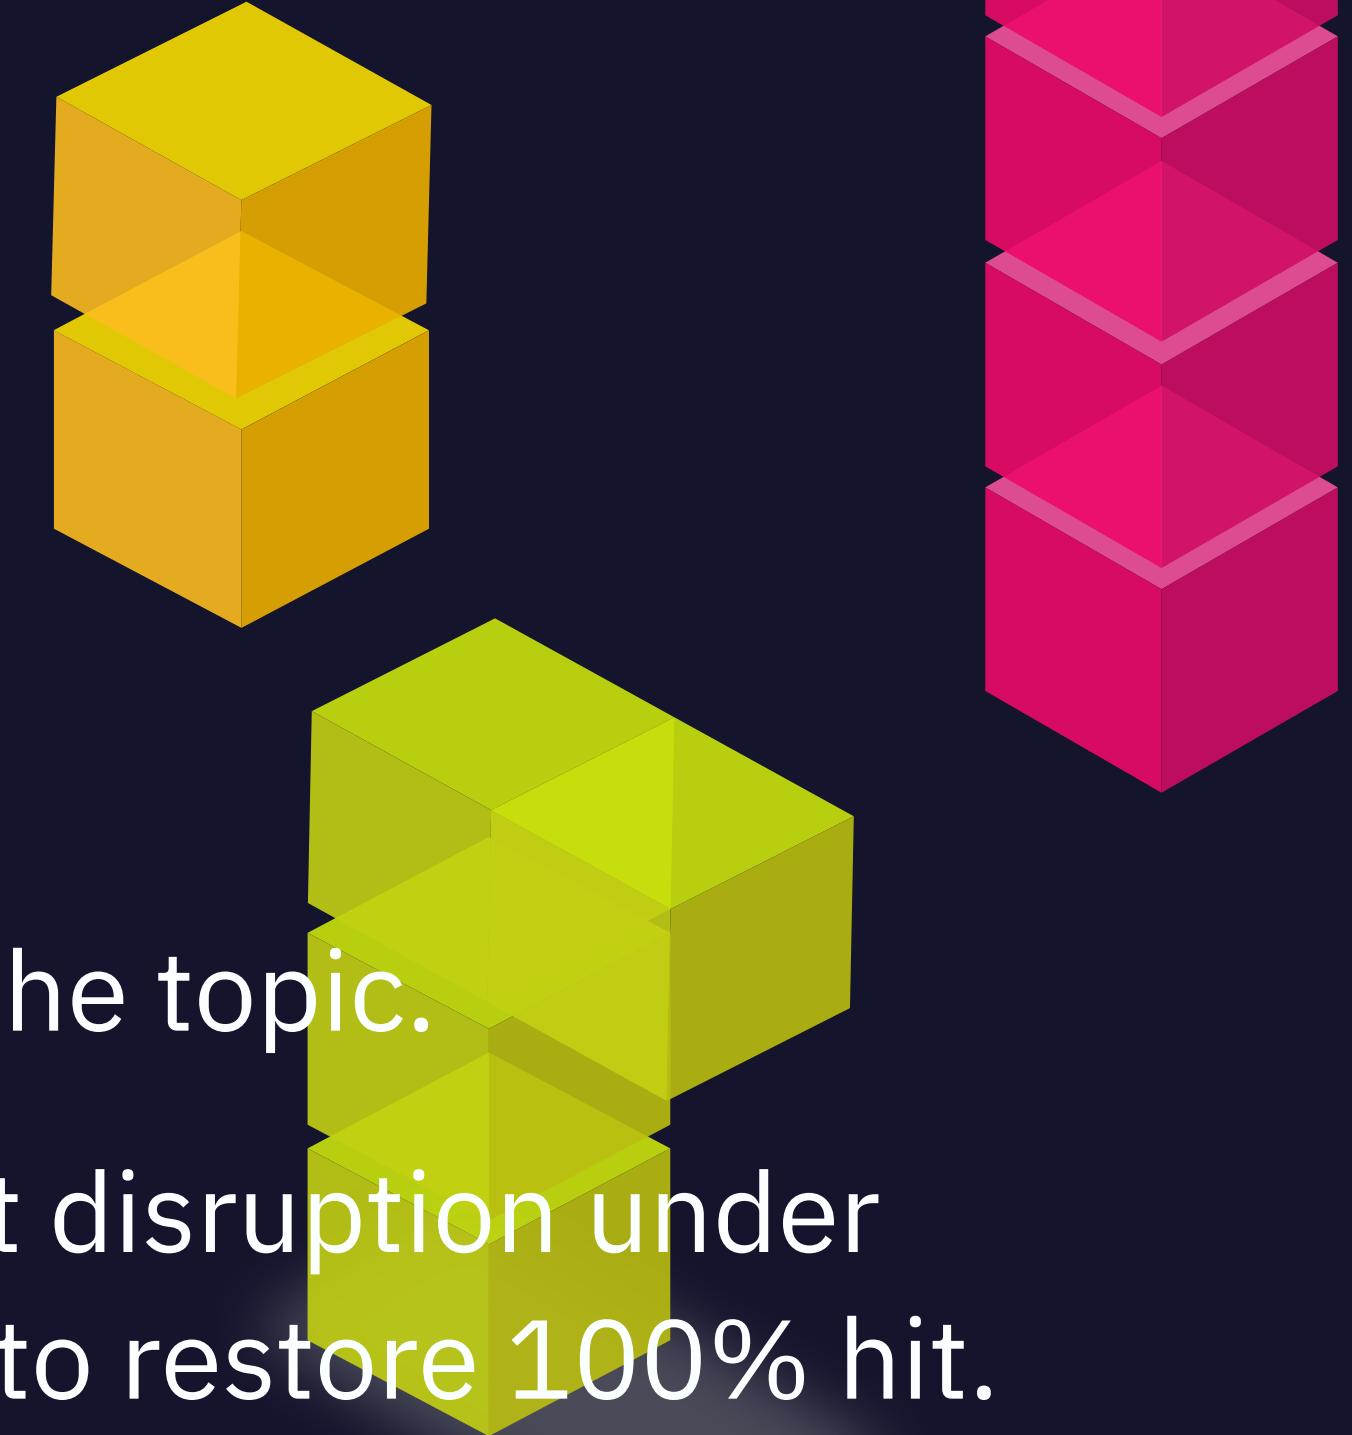
pubsub

PUBSUB / GOSSIP

- Publish/subscribe (pubsub) is a message-oriented communication pattern.
- The interaction model of pubsub is M:N.
 - This is contrast to RPC or request/reply patterns which are 1:1.
- **Pubsub communication is asynchronous and decoupled in nature.**
- Widely popular in enterprise software (Apache ActiveMQ, Apache Kafka, RabbitMQ).
 - **Centralized around brokers in private networks.**
 - **Not suitable for public message fabrics.**



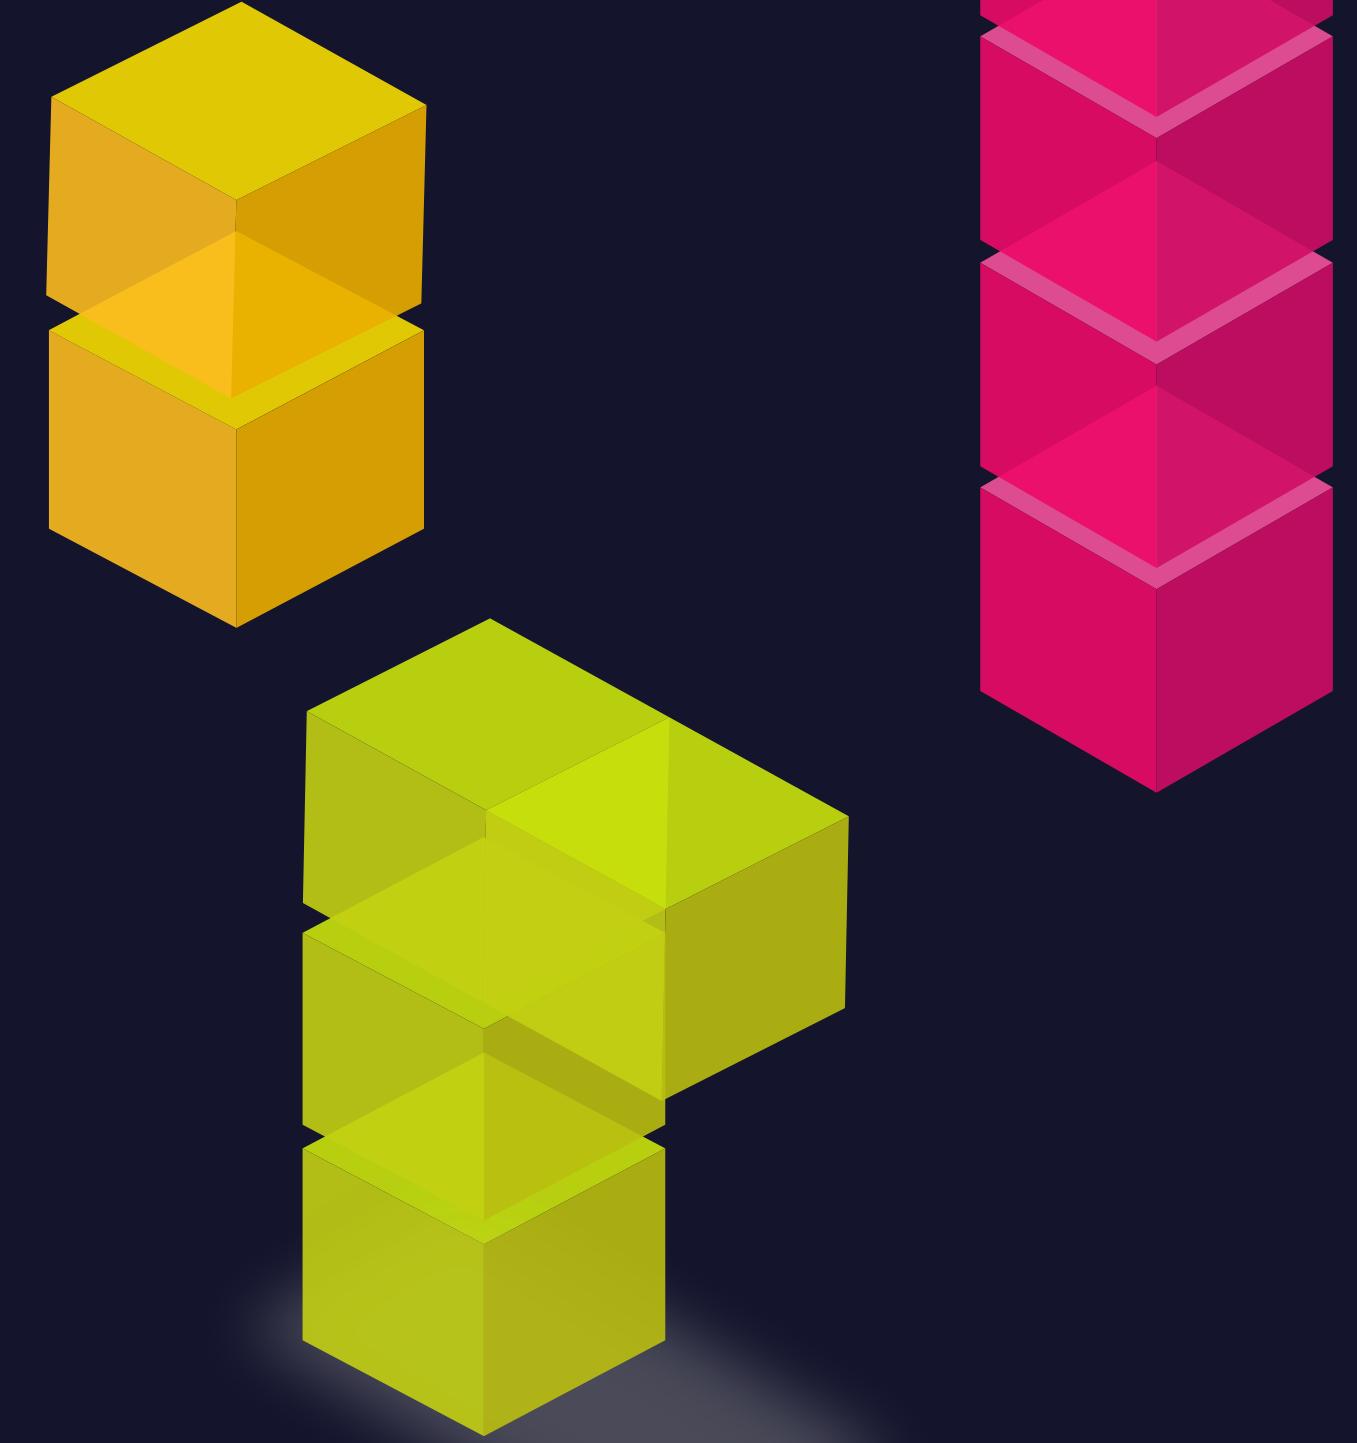
DESIRABLE PROPERTIES IN P2P PUBSUB



1. ***Reliability***: all messages get delivered to all peers subscribed to the topic.
2. ***Resilience***: peers can join and leave topics or the network without disruption under realistic churn scenarios, with **fast recovery** to mend the overlay to restore 100% hit.
3. ***Efficient dissemination***: fast delivery, low write amplification, fair distribution of load.
4. ***Altruism-free routing***: only nodes invested in a topic participate in its dissemination.
5. ***Scalability***: enormous number of nodes, topics, messages, subscriptions, subscribers.
6. ***Low overhead*** of overlay maintenance: nodes track as little state as possible.
7. ***Simplicity***: the system is simple to understand and implement.



DECENTRALIZED PUBSUB IN LIBP2P



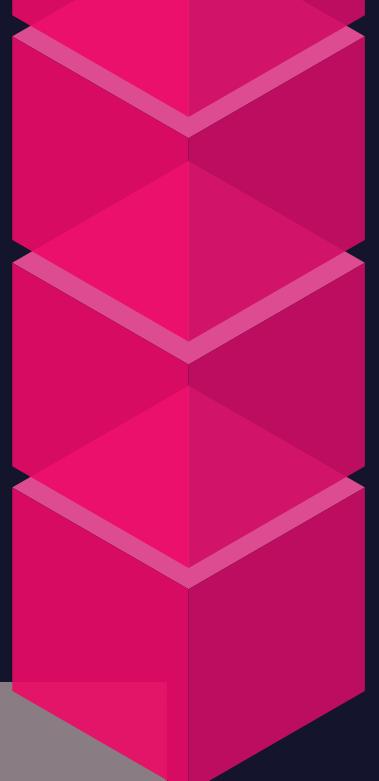
- various routing algorithms available:
 - floodsub (initial) / randomsub (for fun).
 - **gossipsub.**
 - Provides base layer and primitives to build adaptive, self-optimizing routing algorithms capable of repair and rearrangement, in the presence of churn, latency and failures.
 - forms stable reciprocal topic-based meshes to exchange messages.
 - with some gossip to detect message delivery failures and enable routing decisions.
 - control message piggybacking.
 - go, js, rust, jvm implementations



Episub: Plumtree + HyParView + GoCast.

connectivity

CIRCUIT RELAY (TURN)





AUTONAT AND AUTORELAY

- Via the identify protocol, we get decentralised “STUN” for free.
- We also use uPNP and NAT-PMP to create router mappings *automagically*.
- **autonat** is an ambient service that requests dialbacks on addresses we observe of ourselves.
 - it determines if we’re actually dialable and reports a routability: PUBLIC, PRIVATE, UNKNOWN.
- **autorelay** can kick in when routability is PRIVATE.
 - discovers relays automatically, connects to them, and advertises relay addrs. we are now connectable.
 - does not scale well.



next: decentralized hole punching.

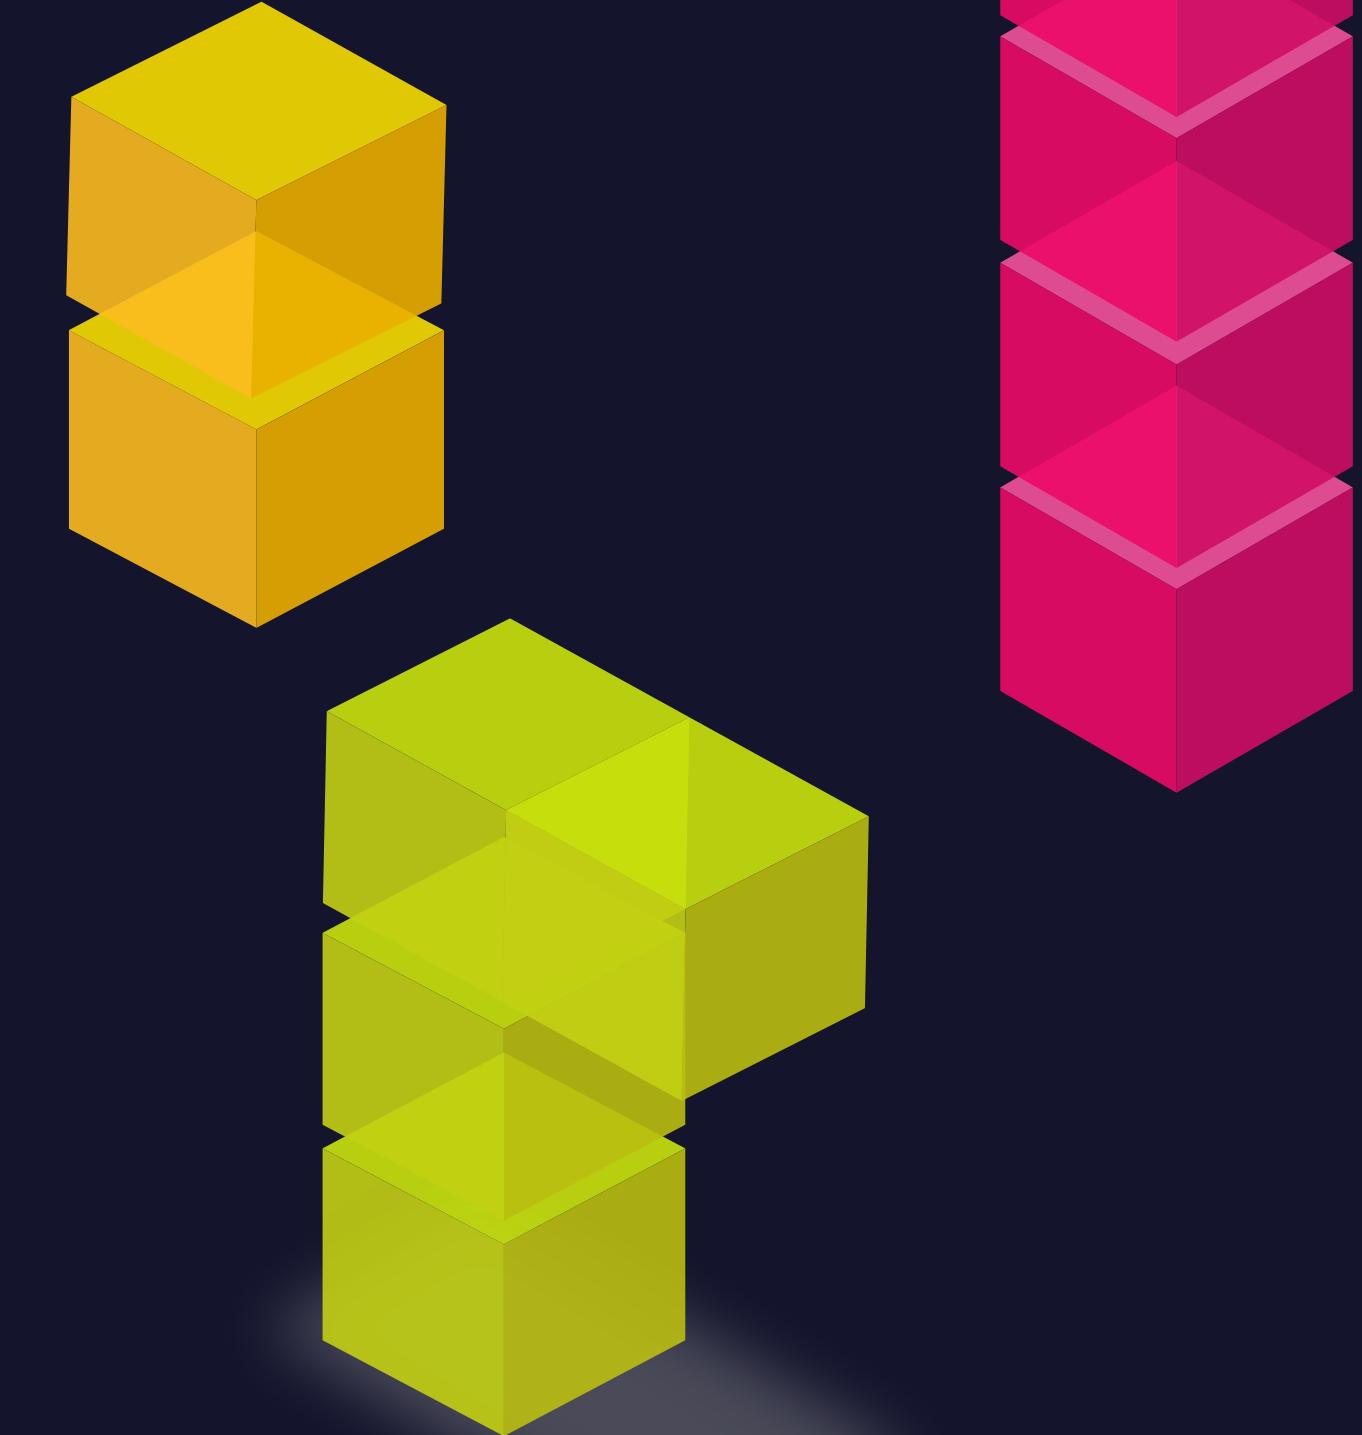




coming soon, to a node near you...

NAT HOLE PUNCHING

- **relaying traffic is inefficient.**
 - latency, centralisation, and expensive.
- **instead: use relay servers as a conduit for hole punching signalling and sync.**
 - establish an embryonic relayed connection.
 - immediately try to upgrade to direct conn
 - then transplant connection state (non-trivial).
 - tcp success rate is 60%; udp is 90%. that's also why we want QUIC.
 - inspired by ICE.
 - careful with tcp simultaneous open; need to adapt multistream-select v1.



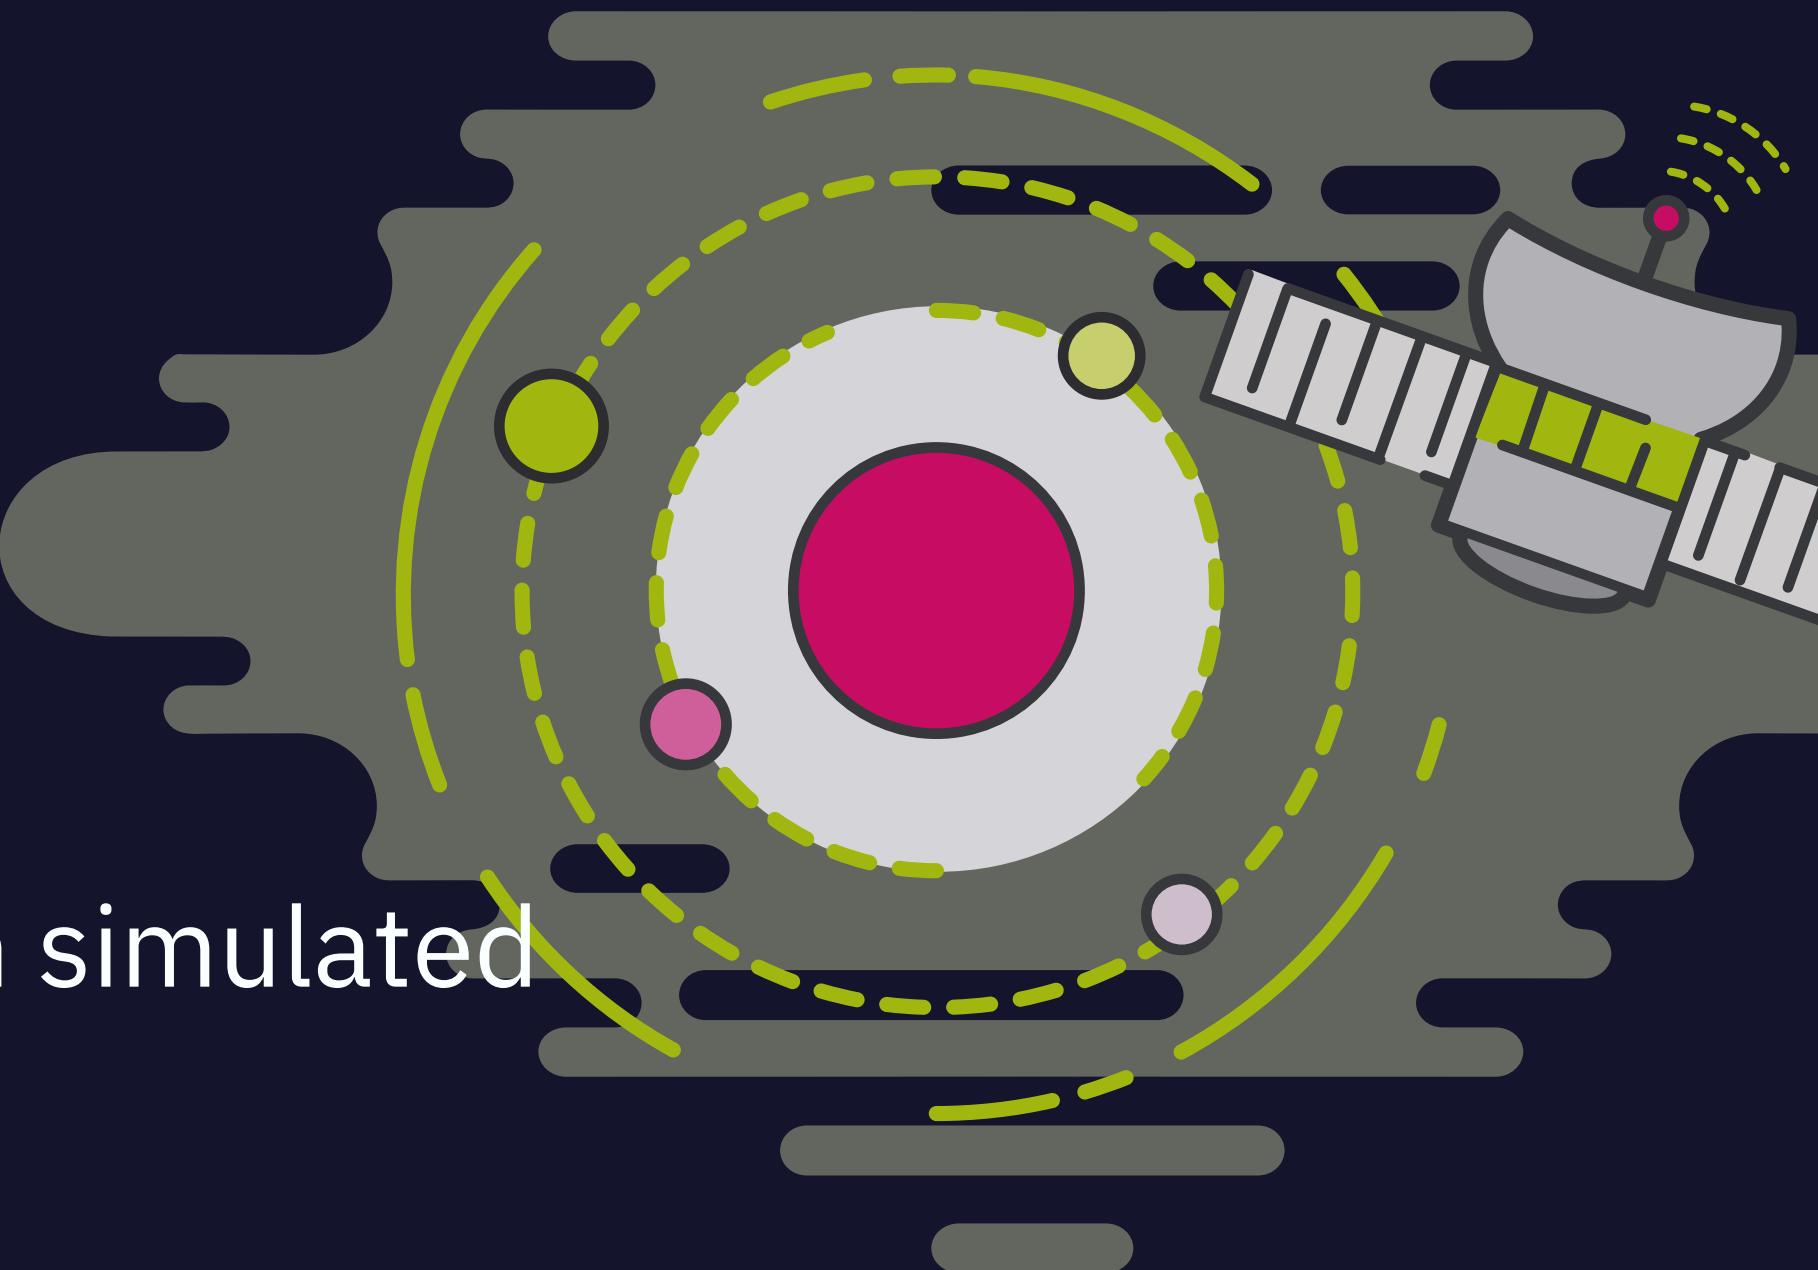


future

(besides everything I've mentioned)

ON THE HORIZON

- distributed, reproduceable testing across 100k nodes with simulated network conditions.
- CTK: conformance test kit.
- even better specs: goal 100% spec'ed.
- Phantom{Drift} (network visualisation).
- DHT 2.0.
- NAT hole punching via relays. Direct connection upgrade.
- episub and gossipsub formalisation.
- spec completeness and accuracy; target: 100%.
- wireshark dissectors.

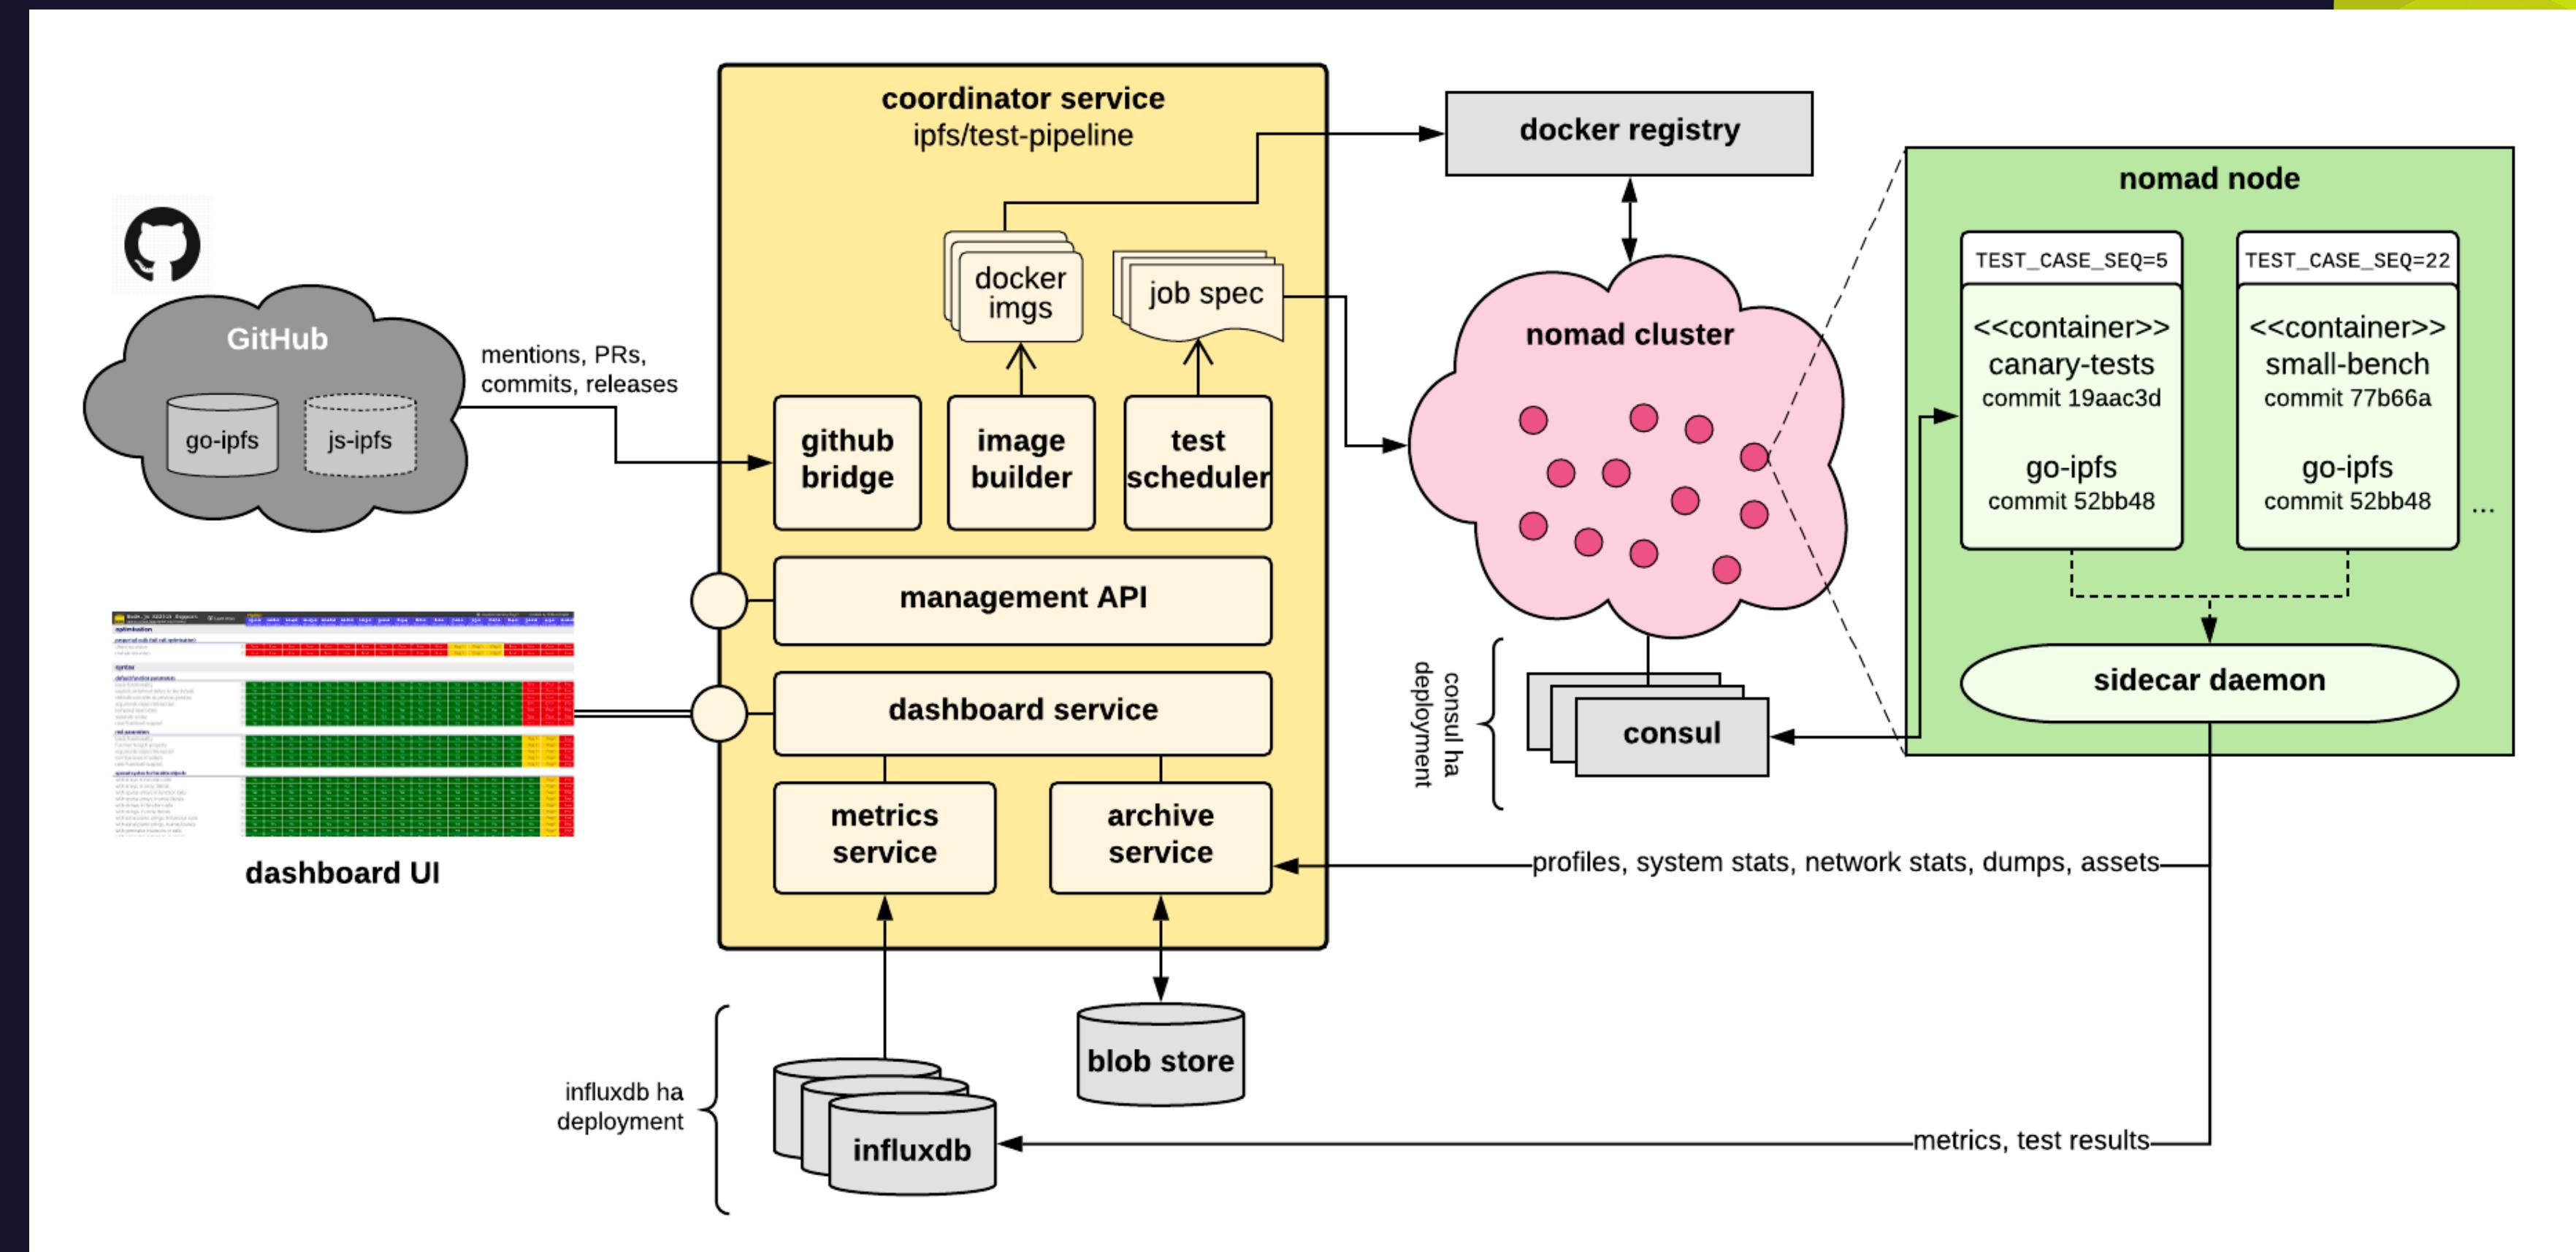




future

on the horizon

INTERPLANETARY TESTGROUND



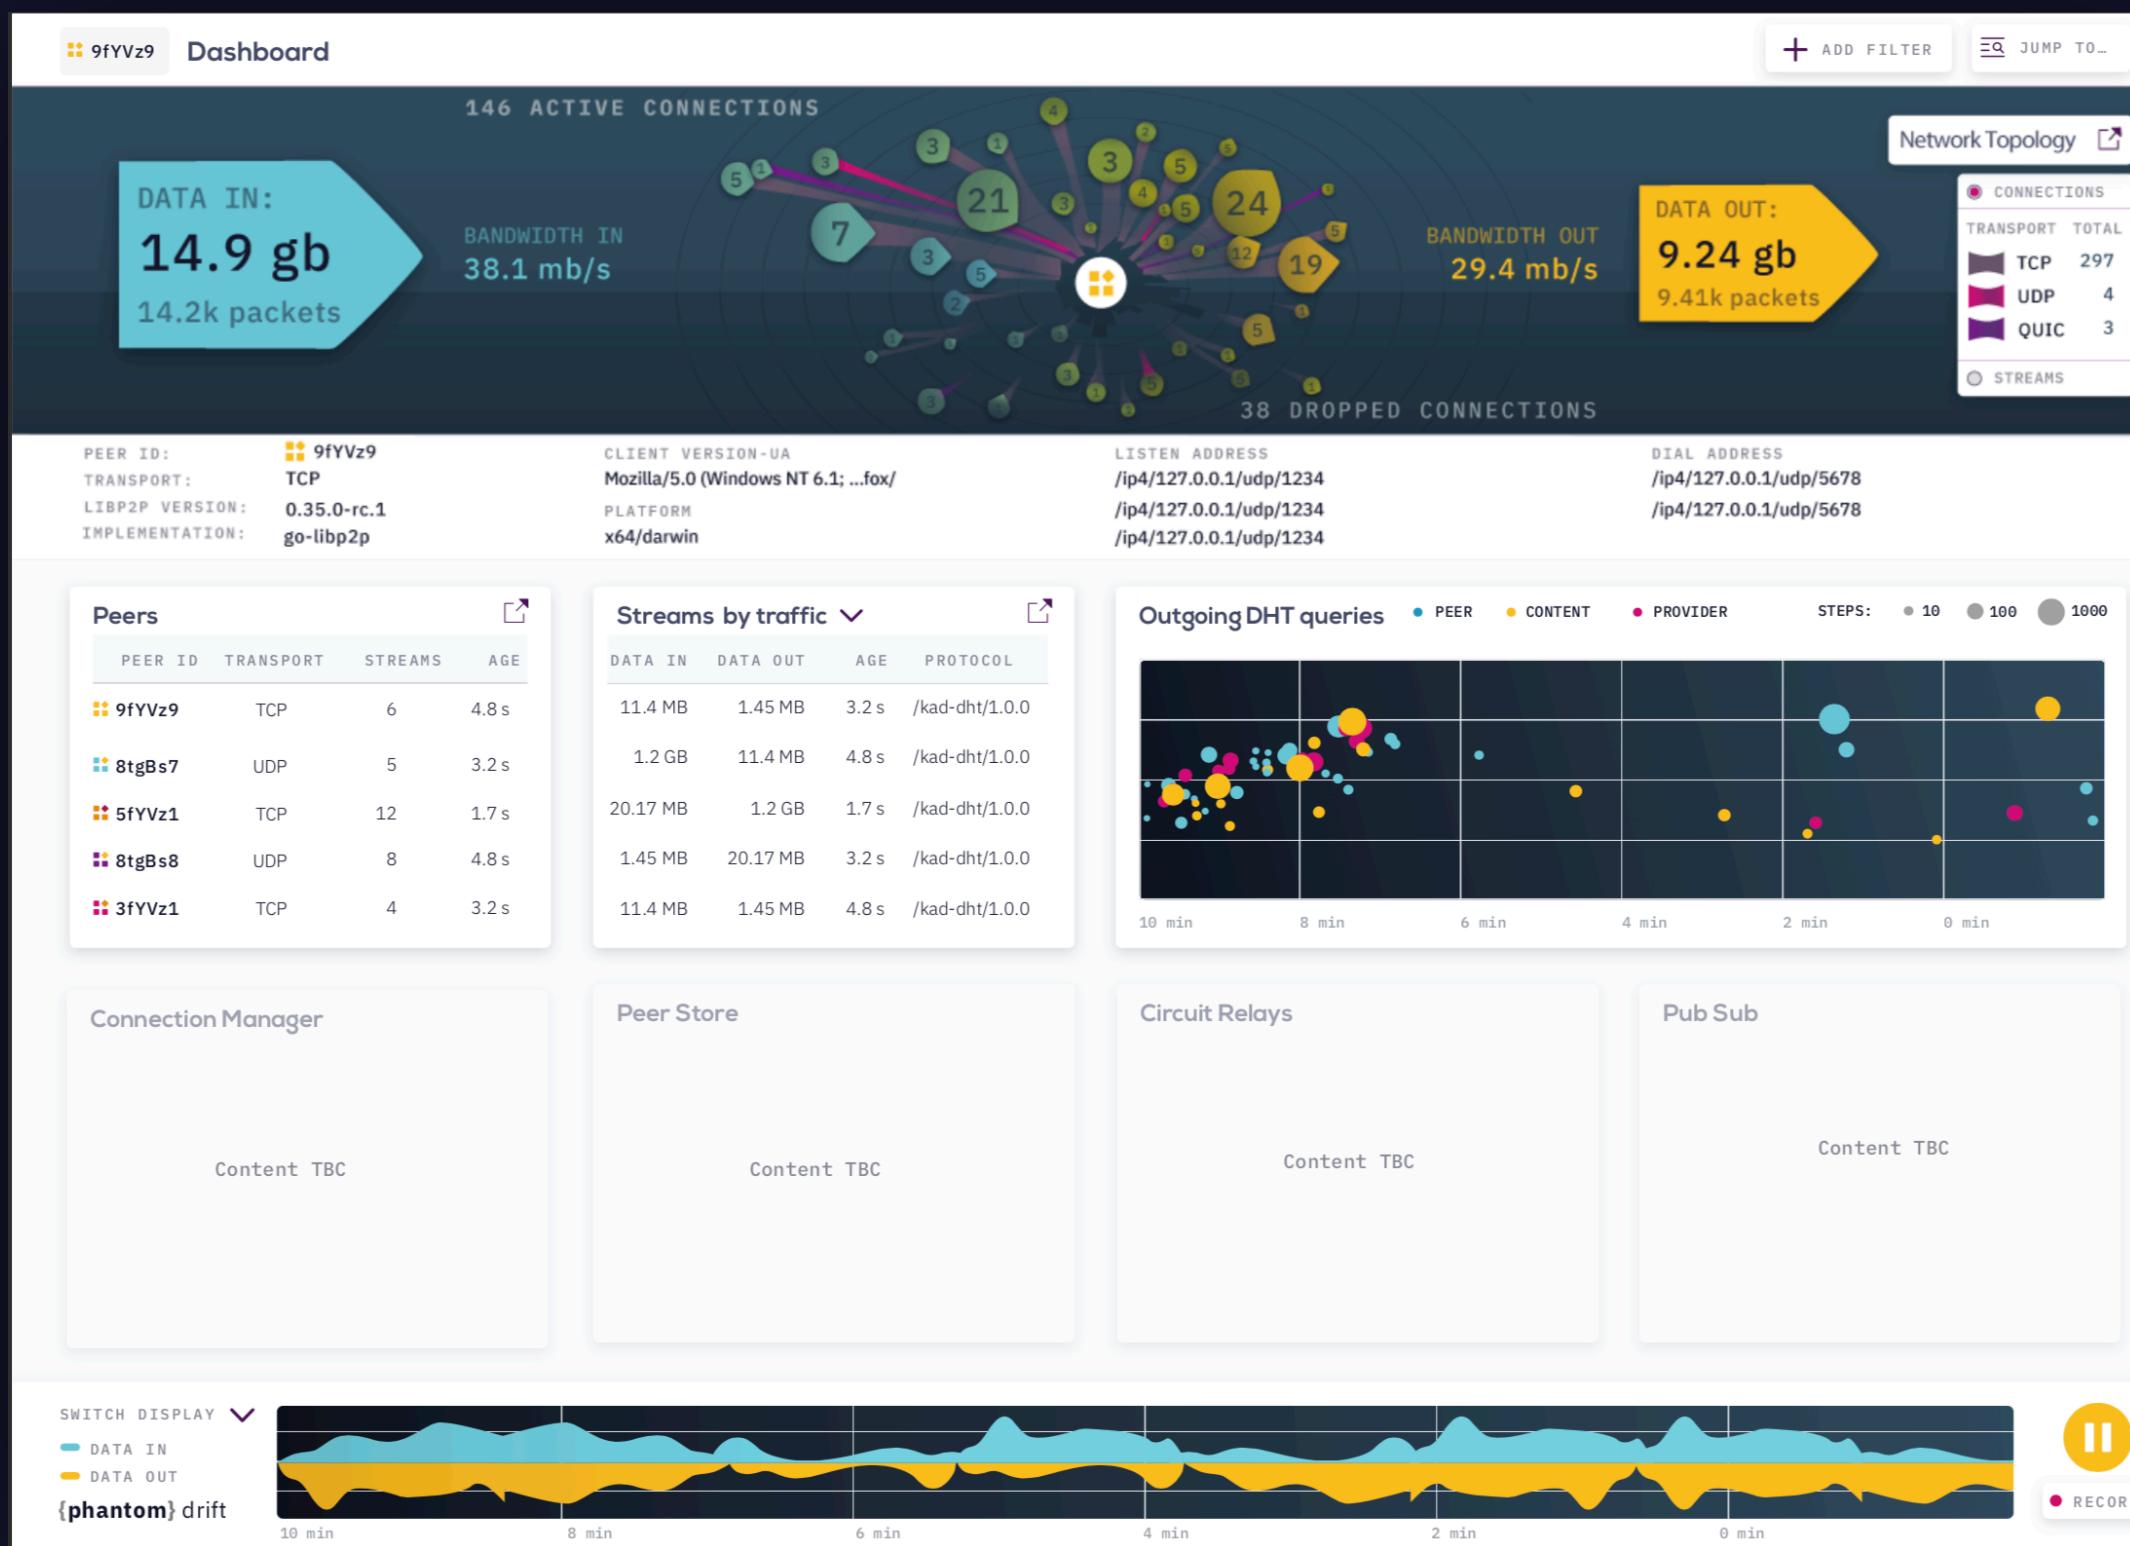
<https://github.com/ipfs/testground>



on the horizon

INSTRUMENTATION

standard introspection protocol and d3-powered visualizations.



WE



YOU

get involved through GitHub!

shape the future



[libp2p/specs](#)

hack on meaty things



[libp2p/devgrants](#)

contribute



project repos

libp2p / devgrants

Unwatch Unstar Fork

Code Issues 15 Pull requests 0 ZenHub Actions Projects 2 Wiki More Settings

want to hack on libp2p? this repo tracks libp2p endeavors eligible for incentivization. [libp2p.io](#)

Edit

bounty libp2p networking Manage topics

28 commits 1 branch 0 releases 2 contributors View license

Branch: master Create new file Find File Clone or download

raulk proposal 005: Conformance Test Kit. (#29)	Latest commit 89f454b 6 days ago
.github/ISSUE_TEMPLATE	Update issue templates
TEMPLATES	readme and template: editorial changes.
004-noise-handshake-implementat...	proposal 004: libp2p Noise handshake implementations. (#21)
005-conformance-test-kit.md	proposal 005: Conformance Test Kit. (#29)
006-gossipsub-threat-proposal.md	proposal 006: refine
008-typescript-bindings.md	proposal 008: TypeScript bindings for js-libp2p-*. (#26)
011-mdns-go-js.md	proposal 011: potential funders += ETHBerlinZwei. (#24)
013-tls13-javascript.md	proposal 013: add Protocol Labs as a potential funder.
016-wireshark-lua-dissectors.md	proposal 016: Wireshark dissectors for libp2p. (#28)
LICENSE-APACHE	flesh out README.md and adjust licenses. (#10)
LICENSE-MIT	flesh out README.md and adjust licenses. (#10)
README.md	readme and template: editorial changes.

<https://github.com/libp2p/devgrants/>



libp2p

questions?

**LET'S HACK ON THE
FUTURE OF P2P
NETWORKING TOGETHER.**

RAÚL KRIPALANI

GITHUB

<https://github.com/raulk>

TWITTER

@raulvk

EMAIL

raul@protocol.ai

